

# CrimeRateAnalysisCode

December 11, 2022

```
[1]: # Import packages

import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import normalize
from matplotlib.pyplot import figure
# from linearmodels import FamaMacBeth
from sklearn.preprocessing import MinMaxScaler
import statsmodels.formula.api as smf
from scipy import stats
import statsmodels.api as sm
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from numpy import mean, std, absolute, sqrt

[2]: # Change this to run in your environment

absolute_file_path = '~/Downloads/DMPProjectRichaRishabh'

[3]: # check our accuracy for each degree, the lower the error the better!

def fetch_best_degree_regression_model(city_data, normalize_original_data,
    ↪for_our_model=True):
    number_degrees = [1,2,3,4,5,6]
    plt_mean_squared_error = []

    if normalize_original_data:
        y_values = city_data['Scaled_CrimeCount']
        if for_our_model:
            x_values = city_data.loc[:,['Gender_ratio', 'Scaled_Unemployment',
            ↪'Scaled_MedianIncomeRate', 'WorkforceCount', 'PercentNegativeUsers',
            ↪'FavorOfDemocrats']]
        else:
```

```

        x_values = city_data.loc[:,['Gender_ratio', 'Scaled_Unemployment',
↪ 'Scaled_MedianIncomeRate']]
    else:
        y_values = city_data['Cul. Crime count']
        if for_our_model:
            x_values = city_data.loc[:,['Gender_ratio', 'UnemploymentRate',
↪ 'MedianIncomeRate', 'WorkforceCount', 'PercentNegativeUsers',
↪ 'FavorOfDemocrats']]
        else:
            x_values = city_data.loc[:,['Gender_ratio', 'UnemploymentRate',
↪ 'MedianIncomeRate']]

    for degree in number_degrees:

        poly_model = PolynomialFeatures(degree=degree)

        poly_x_values = poly_model.fit_transform(x_values)
        poly_model.fit(poly_x_values, y_values)

        regression_model = LinearRegression()
        regression_model.fit(poly_x_values, y_values)
        y_pred = regression_model.predict(poly_x_values)

        plt_mean_squared_error.append(mean_squared_error(y_values, y_pred,
↪ squared=False))

    plt.scatter(number_degrees, plt_mean_squared_error, color="green")
    plt.plot(number_degrees, plt_mean_squared_error, color="red")
    plt.title("Polynomial Regression MSE at different degrees")
    plt.show()

#     return number_degrees[plt_mean_squared_error.
↪ index(min(plt_mean_squared_error))]
    return 2

```

```

[4]: # Plot normalized histogram for model coefficients. coeff_params: dict of
↪ coefficients with values

def plot_hist_for_model_coefficients(coeff_params, p_vals,
↪ min_max_normalization=False, save_fig=False, city_name="", time_period=""):
    i = 0
    coeff_params_significant = {}
    for k, v in coeff_params.items():
        if p_vals[i] <= 0.05:
            coeff_params_significant[k] = v
        i += 1

```

```

    coeff_params_significant = dict(sorted(coeff_params_significant.items(),
↪key=lambda item: item[1], reverse=False))
    if min_max_normalization:
        scaler = MinMaxScaler()
        norm_coeff_values = scaler.fit_transform(pd.
↪DataFrame(coeff_params_significant, index=[0]))
        norm_coeff_values = [i[0] for i in norm_coeff_values]
    else:
        norm_coeff_values = list(dict(coeff_params_significant).values())
        coeff_vars = list(dict(coeff_params_significant).keys())
#     figure(figsize=(10, 8), dpi=80)
    plt.barh(coeff_vars, norm_coeff_values)
    plt.title("Attributes with significant p-values contributing towards crime_
↪rate")
    plt.xlabel("Coefficient Values")
    plt.ylabel("Attributes")
    if save_fig:
        plt.savefig('{_P_VAL_}.png'.format(city_name, time_period),
↪transparent=True, bbox_inches='tight')
    plt.show()

```

```

[5]: def remove_interaction_terms(items_list, indices):
    res = []
    for idx, ele in enumerate(items_list):
        if idx not in indices:
            res.append(ele)
    return res

```

```

[6]: # Removing interaction terms with less p-val in LA
# Parameters: x=transformed x, z=interaction variable names
# Removing following variables from Polynomial Equation
# 3: WorkforceCount
# 7: Gender_ratio^2
# 10: Gender_ratio WorkforceCount
# 11: Gender_ratio PercentNegativeUsers
# 12: Gender_ratio FavorOfDemocrats
# 13: Gender_ratio SearchCountForDepression
# 14: UnemploymentRate^2
# 15: UnemploymentRate MedianIncomeRate
# 17: UnemploymentRate PercentNegativeUsers
# 18: UnemploymentRate FavorOfDemocrats
# 19: UnemploymentRate SearchCountForDepression
# 20: MedianIncomeRate^2
# 22: MedianIncomeRate PercentNegativeUsers
# 23: MedianIncomeRate FavorOfDemocrats
# 24: MedianIncomeRate SearchCountForDepression
# 25: WorkforceCount^2

```

```

# 28: WorkforceCount SearchCountForDepression
# 30: PercentNegativeUsers FavorOfDemocrats
# 31: PercentNegativeUsers SearchCountForDepression

def remove_interaction_terms_for_LA(x, z):
    new_x_ = []
    interaction_terms_indices = [3, 7, 10, 11, 12, 13, 14, 15, 17, 18, 19, 20, 22, 23, 24, 25, 28, 30, 31]
    for item in x:
        temp = remove_interaction_terms(item, interaction_terms_indices)
        new_x_.append(temp)

    new_x_ = pd.DataFrame(np.array(new_x_), columns=remove_interaction_terms(z, interaction_terms_indices))
    return new_x_

```

```

[7]: # Removing interaction terms with less p-val in NYC
# Parameters: x=transformed x, z=interaction variable names
# Removing following variables from Polynomial Equation
# 1: UnemploymentRate
# 6: SearchCountForDepression
# 11: Gender_ratio PercentNegativeUsers
# 13: Gender_ratio SearchCountForDepression
# 14: UnemploymentRate^2
# 17: UnemploymentRate PercentNegativeUsers
# 19: UnemploymentRate SearchCountForDepression
# 21: MedianIncomeRate WorkforceCount
# 22: MedianIncomeRate PercentNegativeUsers
# 24: MedianIncomeRate SearchCountForDepression
# 26: WorkforceCount PercentNegativeUsers
# 28: WorkforceCount SearchCountForDepression
# 30: PercentNegativeUsers FavorOfDemocrats
# 31: PercentNegativeUsers SearchCountForDepression
# 33: FavorOfDemocrats SearchCountForDepression
# 34: SearchCountForDepression^2

def remove_interaction_terms_for_NYC(x, z):
    new_x_ = []
    interaction_terms_indices = [1, 6, 11, 13, 14, 17, 19, 21, 22, 24, 26, 28, 30, 31, 33, 34]
    for item in x:
        temp = remove_interaction_terms(item, interaction_terms_indices)
        new_x_.append(temp)

    new_x_ = pd.DataFrame(np.array(new_x_), columns=remove_interaction_terms(z, interaction_terms_indices))
    return new_x_

```

```
[8]: # Removing interaction terms with less p-val in CHICAGO
# Parameters: x=transformed x, z=interaction variable names

def remove_interaction_terms_for_CHICAGO(x, z):
    new_x_ = []
    interaction_terms_indices = [1, 8, 14, 17, 19, 22, 24, 26, 28, 29, 30, 31,
↪33]
    for item in x:
        temp = remove_interaction_terms(item, interaction_terms_indices)
        new_x_.append(temp)

    new_x_ = pd.DataFrame(np.array(new_x_), columns=remove_interaction_terms(z,
↪interaction_terms_indices))
    return new_x_
```

```
[9]: # Performs 10-cross validation on our dataset to find RMSE

def perform_k_cross_validation(X, y):
    #define cross-validation method to use
    cv = KFold(n_splits=10, random_state=1, shuffle=True)

    #build multiple linear regression model
    model = LinearRegression()

    #use k-fold CV to evaluate model
    scores = cross_val_score(model, X, y, scoring='neg_mean_squared_error',
                             cv=cv, n_jobs=-1)

    #return MSE, std
    return mean(absolute(scores)), std(absolute(scores))
```

```
[10]: # Ground Truth

mse_ground_truth_before_covid, std_ground_truth_before_covid = {}, {}
mse_ground_truth, std_ground_truth = {}, {}

def regress_ground_truth(city_name, city_data, best_degree,
↪normalize_original_data, poly_model, before_covid=False):
    lin_model_p_vals, poly_model_p_vals = [], []
    if normalize_original_data:
        y = city_data['Scaled_CrimeCount']
        x = city_data.loc[:, ['Gender_ratio', 'Scaled_Unemployment',
↪'Scaled_MedianIncomeRate']]
    else:
        y = city_data['Cul. Crime count']
        x = city_data.loc[:, ['Gender_ratio', 'UnemploymentRate',
↪'MedianIncomeRate']]
```

```

X = sm.add_constant(x)
lin_model = sm.OLS(y, X).fit()
for attributeIndex in range(0, len(X.columns)):
    lin_model_p_vals.append(lin_model.pvalues[attributeIndex])
lin_model_summary = lin_model.summary()
if before_covid:
    mse_ground_truth_before_covid[city_name],
std_ground_truth_before_covid[city_name] = perform_k_cross_validation(X, y)
else:
    mse_ground_truth[city_name], std_ground_truth[city_name] =
perform_k_cross_validation(X, y)
# linear_model = LinearRegression()
# linear_model.fit(x, y)
# r_sq_lm = linear_model.score(x, y)
# features = ['Intercept'] + ['Gender_ratio', 'Scaled_Unemployment',
'Scaled_MedianIncomeRate']
# p_val_lm = get_p_val(linear_model, x, y, features)
if poly_model:
    transformer = PolynomialFeatures(degree=best_degree, include_bias=False)
    x_ = transformer.fit_transform(x)
    x_ = pd.DataFrame(np.array(x_), columns=transformer.
get_feature_names_out())
    X = sm.add_constant(x_)
    poly_model = sm.OLS(list(y), X).fit()
    for attributeIndex in range(0, len(X.columns)):
        poly_model_p_vals.append(poly_model.pvalues[attributeIndex])
    poly_model_summary = poly_model.summary()
# poly_model = LinearRegression().fit(x_, y)
# r_sq_pm = poly_model.score(x_, y)
# features = ['Intercept'] + list(transformer.get_feature_names_out())
# p_val_pm = get_p_val(poly_model, x_, y, features)

print("Ground Truth:\nLinear Model summary:\n {}".format(lin_model_summary))
# print("Model Coefficients: \nLinear: {}".format(list(linear_model.
coef_)))
print("\n\nPlot for Linear Model Coefficients for Ground Truth:\n")
plot_hist_for_model_coefficients(lin_model.params, lin_model_p_vals)
# print("Polynomial: {}".format(list(poly_model.coef_)))
if poly_model:
    print("\n\nPolynomial Model (degree={}) summary:\n {}".format(
best_degree, poly_model_summary))
    print("\n\nPlot for Polynormial Model Coefficients:\n")
    plot_hist_for_model_coefficients(poly_model.params, poly_model_p_vals)

```

```

[11]: # Our Model

mse_our_model_before_covid, std_our_model_before_covid = {}, {}
mse_our_model, std_our_model = {}, {}

def regress_our_model(city_name, city_data, best_degree,
    ↪normalize_original_data, include_mental_health_data, poly_model,
    ↪covid=False, before_covid=False, save_fig=False, time_period=""):
    lin_model_p_vals, poly_model_p_vals = [], []
    if normalize_original_data:
        y = city_data['Scaled_CrimeCount']
        if include_mental_health_data:
            ori_features = ['Gender_ratio', 'Scaled_Unemployment',
    ↪'Scaled_MedianIncomeRate', 'WorkforceCount', 'PercentNegativeUsers',
    ↪'FavorOfDemocrats', 'Scaled_SearchCountForDepression']
        else:
            ori_features = ['Gender_ratio', 'Scaled_Unemployment',
    ↪'Scaled_MedianIncomeRate', 'WorkforceCount', 'PercentNegativeUsers',
    ↪'FavorOfDemocrats']
        elif include_mental_health_data:
            y = city_data['Cul. Crime count']
            ori_features = ['Gender_ratio', 'UnemploymentRate', 'MedianIncomeRate',
    ↪'WorkforceCount', 'PercentNegativeUsers', 'FavorOfDemocrats',
    ↪'SearchCountForDepression']
        else:
            y = city_data['Cul. Crime count']
            ori_features = ['Gender_ratio', 'UnemploymentRate', 'MedianIncomeRate',
    ↪'WorkforceCount', 'PercentNegativeUsers', 'FavorOfDemocrats']
    if covid:
        ori_features += ['CovidPercPositive']

    x = city_data.loc[:,ori_features]

    X = sm.add_constant(x)
    lin_model = sm.OLS(y, X).fit()
    for attributeIndex in range (0, len(X.columns)):
        lin_model_p_vals.append(lin_model.pvalues[attributeIndex])
    lin_model_summary = lin_model.summary()
    if not covid:
        if before_covid:
            mse_our_model_before_covid[city_name],
    ↪std_our_model_before_covid[city_name] = perform_k_cross_validation(X, y)
        else:
            mse_our_model[city_name], std_our_model[city_name] =
    ↪perform_k_cross_validation(X, y)
    # linear_model = LinearRegression()

```

```

# linear_model.fit(x, y)
# r_sq_lm = linear_model.score(x, y)
# features = ['Intercept'] + ori_features
# p_val_lm = get_p_val(linear_model, x, y, features)

# poly = PolynomialFeatures(degree = 2)
# X_poly = poly.fit_transform(x)
# poly.fit(X_poly, y)
# poly_model = LinearRegression()
# poly_model.fit(X_poly, y)
# r_sq_pm = poly_model.score(X_poly, y)
if poly_model:
    transformer = PolynomialFeatures(degree=best_degree, include_bias=False)
    x_ = transformer.fit_transform(x)
    z = transformer.get_feature_names_out()
    if city_name == "LA" and not covid:
        x_ = remove_interaction_terms_for_LA(x_, z)
    elif city_name == "NYC" and not covid:
        x_ = remove_interaction_terms_for_NYC(x_, z)
    elif city_name == "CHICAGO" and not covid:
        x_ = remove_interaction_terms_for_CHICAGO(x_, z)
    else:
        x_ = pd.DataFrame(np.array(x_), columns=z)
    X = sm.add_constant(x_)
    poly_model = sm.OLS(list(y), X).fit()
    for attributeIndex in range(0, len(X.columns)):
        poly_model_p_vals.append(poly_model.pvalues[attributeIndex])
    poly_model_summary = poly_model.summary()
# poly_model = LinearRegression().fit(x_, y)
# r_sq_pm = poly_model.score(x_, y)
# features = ['Intercept'] + list(transformer.get_feature_names_out())
# p_val_pm = get_p_val(poly_model, x_, y, features)

print("Our Model:\nLinear Model summary:\n {}".format(lin_model_summary))
# print("Model Coefficients: \nLinear: {}".format(list(linear_model.
    ↪coef_)))
print("\n\nPlot for Linear Model Coefficients for Our Model:\n")
plot_hist_for_model_coefficients(lin_model.params, lin_model_p_vals)
# print("Polynomial: {}".format(list(poly_model.coef_)))
if poly_model:
    print("\n\nPolynomial Model (degree={}) summary:\n {}".format(
    ↪best_degree, poly_model_summary))
    print("\n\nPlot for Polynormial Model Coefficients:\n")
    plot_hist_for_model_coefficients(poly_model.params, poly_model_p_vals)

```

[12]: # Our Model without ground truth columns



```

# mse_our_model_before_covid, std_our_model_before_covid = {}, {}
# mse_our_model, std_our_model = {}, {}

def regress_our_model_without_ground_truth(city_name, city_data,
    ↪normalize_original_data, include_mental_health_data, poly_model,
    ↪covid=False, before_covid=False, save_fig=False, time_period=""):
    lin_model_p_vals, poly_model_p_vals = [], []
    if normalize_original_data:
        y = city_data['Scaled_CrimeCount']
        if include_mental_health_data:
            ori_features = ['WorkforceCount', 'PercentNegativeUsers',
    ↪'FavorOfDemocrats', 'Scaled_SearchCountForDepression']
        else:
            ori_features = ['WorkforceCount', 'PercentNegativeUsers',
    ↪'FavorOfDemocrats']
        elif include_mental_health_data:
            y = city_data['Cul. Crime count']
            ori_features = ['WorkforceCount', 'PercentNegativeUsers',
    ↪'FavorOfDemocrats', 'SearchCountForDepression']
        else:
            y = city_data['Cul. Crime count']
            ori_features = ['WorkforceCount', 'PercentNegativeUsers',
    ↪'FavorOfDemocrats']
    if covid:
        ori_features += ['CovidPercPositive']

    x = city_data.loc[:,ori_features]

    X = sm.add_constant(x)
    lin_model = sm.OLS(y, X).fit()
    for attributeIndex in range(0, len(X.columns)):
        lin_model_p_vals.append(lin_model.pvalues[attributeIndex])
    lin_model_summary = lin_model.summary()
    #     if not covid:
    #         if before_covid:
    #             mse_our_model_before_covid[city_name],
    ↪std_our_model_before_covid[city_name] = perform_k_cross_validation(X, y)
    #         else:
    #             mse_our_model[city_name], std_our_model[city_name] =
    ↪perform_k_cross_validation(X, y)
    #     linear_model = LinearRegression()
    #     linear_model.fit(x, y)
    #     r_sq_lm = linear_model.score(x, y)
    #     features = ['Intercept'] + ori_features
    #     p_val_lm = get_p_val(linear_model, x, y, features)

```

```

# poly = PolynomialFeatures(degree = 2)
# X_poly = poly.fit_transform(x)
# poly.fit(X_poly, y)
# poly_model = LinearRegression()
# poly_model.fit(X_poly, y)
# r_sq_pm = poly_model.score(X_poly, y)
if poly_model:
    transformer = PolynomialFeatures(degree=best_degree, include_bias=False)
    x_ = transformer.fit_transform(x)
    z = transformer.get_feature_names_out()
    if city_name == "LA" and not covid:
        x_ = remove_interaction_terms_for_LA(x_, z)
    elif city_name == "NYC" and not covid:
        x_ = remove_interaction_terms_for_NYC(x_, z)
    elif city_name == "CHICAGO" and not covid:
        x_ = remove_interaction_terms_for_CHICAGO(x_, z)
    else:
        x_ = pd.DataFrame(np.array(x_), columns=z)
    X = sm.add_constant(x_)
    poly_model = sm.OLS(list(y), X).fit()
    for attributeIndex in range(0, len(X.columns)):
        poly_model_p_vals.append(poly_model.pvalues[attributeIndex])
    poly_model_summary = poly_model.summary()
# poly_model = LinearRegression().fit(x_, y)
# r_sq_pm = poly_model.score(x_, y)
# features = ['Intercept'] + list(transformer.get_feature_names_out())
# p_val_pm = get_p_val(poly_model, x_, y, features)

print("Our Model without ground truth:\nLinear Model summary:\n {}".format(
lin_model_summary))
# print("Model Coefficients: \nLinear: {}".format(list(linear_model.coef_)))
print("\n\nPlot for Linear Model Coefficients for Our Model without ground_
truth:\n")
plot_hist_for_model_coefficients(lin_model.params, lin_model_p_vals)
# print("Polynomial: {}".format(list(poly_model.coef_)))
if poly_model:
    print("\n\nPolynomial Model (degree={}) summary:\n {}".format(
best_degree, poly_model_summary))
    print("\n\nPlot for Polynomial Model Coefficients:\n")
    plot_hist_for_model_coefficients(poly_model.params, poly_model_p_vals)

```

[13]: # Plot Correlation graphs

```

def plot_correlation_graphs(city_data, normalize_original_data):
    print("\nCorrelation Graphs:\n")

```

```

if normalize_original_data:
    unemployment_col = city_data['Scaled_Unemployment']
    median_income_col = city_data['Scaled_MedianIncomeRate']
    y = city_data['Scaled_CrimeCount']
else:
    unemployment_col = city_data['UnemploymentRate']
    median_income_col = city_data['MedianIncomeRate']
    y = city_data['Cul. Crime count']

figure(figsize=(15, 10), dpi=80)

plt.subplot(2, 3, 1)
plt.scatter(city_data['Gender_ratio'], y)
plt.title("Gender Ratio v/s Crime Count")

plt.subplot(2, 3, 2)
plt.scatter(unemployment_col, y)
plt.title("Scaled_Unemployment v/s Crime Count")

plt.subplot(2, 3, 3)
plt.scatter(median_income_col, y)
plt.title("Scaled_MedianIncomeRate v/s Crime Count")

plt.subplot(2, 3, 4)
plt.scatter(city_data['WorkforceCount'], y)
plt.title("WorkforceCount v/s Crime Count")

plt.subplot(2, 3, 5)
plt.scatter(city_data['PercentNegativeUsers'], y)
plt.title("PercentNegativeUsers v/s Crime Count")

plt.subplot(2, 3, 6)
plt.scatter(city_data['FavorOfDemocrats'], y)
plt.title("FavorOfDemocrats v/s Crime Count")

plt.show()

```

```
[14]: # Append multiple city dataframes into single list
```

```

# def build_single_dataset(city_name, city_data, all_cities_df):
#     city_data['City'] = city_name
#     all_cities_df.append(city_data)

```

```
[15]: # Min-Max Normalization on full data
```

```

def normalize_dataset(city_data, include_mental_health_data):
    scaler = MinMaxScaler()

```

```

    city_data[["Scaled_CrimeCount"]] = scaler.fit_transform(city_data[['Cul.
↳Crime count']])
    city_data[["Scaled_Unemployment"]] = scaler.
↳fit_transform(city_data[['UnemploymentRate']])
    city_data[["Scaled_MedianIncomeRate"]] = scaler.
↳fit_transform(city_data[['MedianIncomeRate']])
    if include_mental_health_data:
        city_data[["Scaled_SearchCountForDepression"]] = scaler.
↳fit_transform(city_data[['SearchCountForDepression']])

```

```

[16]: # Covid File Name Structure - <CityNameInCaps>_Covid.csv, No Header, First
↳Column as Percent and Second Column as MonthYear

def infer_covid_comparison(city_name, city_data, absolute_file_path,
↳best_degree, normalize_original_data, include_mental_health_data,
↳poly_model):
    city_covid = pd.read_csv("{}_{}/{}_Covid.csv".format(absolute_file_path,
↳city_name), header=None)
    city_covid.columns = ['Perc', 'MonthYear']
    city_covid['MonthYear'] = pd.to_datetime(city_covid['MonthYear'],
↳format="%b-%Y")
    city_covid = city_covid.sort_values(by=['MonthYear'])

    city_data['MonthYear'] = pd.to_datetime(city_data['MonthYear'],
↳format="%b-%Y")

    city_combined_dataset_covid = city_data[city_data['MonthYear'] >=
↳'2020-03-01']

    perc = []
    for index, row in city_combined_dataset_covid.iterrows():
        month_year = row['MonthYear']
        perc_pos_covid_cases = city_covid[city_covid['MonthYear'] ==
↳month_year]['Perc']
        perc.append(float(perc_pos_covid_cases))

    city_combined_dataset_covid['CovidPercPositive'] = perc

    print("\nCovid Data Comparison:\n\nRegression from Mar-2020 to Dec-2021,
↳along with Covid Positive Cases Per Month:\n")
    regress_our_model(city_name, city_combined_dataset_covid, best_degree,
↳normalize_original_data, include_mental_health_data, poly_model, covid=True)
    print("\nRegression for the same time frame without considering Covid Cases:
↳\n")
    regress_our_model(city_name, city_combined_dataset_covid, best_degree,
↳normalize_original_data, include_mental_health_data, poly_model)

```

```
[17]: # File Name Structure - <CityNameInCaps>_GoogleTrends.csv, with original file
      ↪format as downloaded from google trends
```

```
def add_mental_health_data(city_name, city_data, absolute_file_path):
    mh_data = pd.read_csv("{}_GoogleTrends.csv".format(absolute_file_path,
    ↪city_name), header=1)
    mh_data['Month'] = pd.to_datetime(mh_data['Month'], format="%Y-%m")
    mh_data = mh_data.sort_values(by=['Month'])
    mh_data.columns = ['Month', 'Count']

    city_data['MonthYear'] = pd.to_datetime(city_data['MonthYear'],
    ↪format="%b-%y")
    city_data = city_data.sort_values(by=['MonthYear'])
    city_data['SearchCountForDepression'] = list(mh_data['Count'])
    # city_data.to_csv("{}-withMentalData.csv".format(city_name))
    return city_data
```

```
[18]: # For few cities, MonthYear Column was missing, so created on the fly
```

```
def create_month_year_column(city_data):
    city_data['YearMonth'] = pd.to_datetime(city_data['YearMonth'],
    ↪format="%y-%b")
    city_data["MonthYear"] = city_data['YearMonth'].dt.strftime("%b-%y")
    return city_data
```

```
[19]: def fetch_city_data_before_covid(city_data):
      city_data['Time'] = pd.to_datetime(city_data['MonthYear'], format="%b-%y")
      return city_data[city_data['Time'] < '2020-03-01']
```

```
[20]: def plot_time_series_graph(city_name, city_data, save_fig=False):
      plt.figure(figsize=(16, 8), dpi=150)
      # city_data['WorkforceCount'].plot(label='Workforce%', color='orange')
      # city_data['Scaled_CrimeCount'].plot(label='Crime%')
      # plt.title('{}-Workforce% v/s Crime%'.format(city_name))
      # plt.xlabel('Time Period (Months) - 2011-2021')
      # plt.legend()
      fig, ax1 = plt.subplots()

      ax2 = ax1.twinx()
      ax1.plot(city_data["MonthYear"], city_data['Scaled_Unemployment'], 'g-',
      ↪color="#6495ED")
      ax2.plot(city_data["MonthYear"], city_data['Scaled_CrimeCount'], 'b-',
      ↪color="#CC4F1B")

      ax1.set_xlabel('Time Period (Months) - 2011-2021')
      ax1.set_ylabel('Unemployment%', color='#6495ED')
```

```

ax1.set_title('{}-Unemployment% v/s Crime%'.format(city_name))
ax2.set_ylabel('Crime%', color='#CC4F1B')

# plt.savefig('{}_TIME_SERIES_Unemployment.png'.format(city_name),
↳transparent=True, bbox_inches='tight')
plt.show()

```

```

[21]: # Driver Function to read and process data for all cities

def infer_data(city_name, city_file, all_cities_df, absolute_file_path,
↳normalize_original_data=True, include_mental_health_data=True,
↳poly_model=True):

    city_data = pd.read_csv("{} / {}".format(absolute_file_path, city_file),
↳header=0)

    if city_name in ["SEATTLE", "DENVER", "DALLAS", "NEW_ORLEANS", "NYC"]:
        city_data = create_month_year_column(city_data)

    if include_mental_health_data:
        city_data = add_mental_health_data(city_name, city_data,
↳absolute_file_path)
    if normalize_original_data:
        normalize_dataset(city_data, include_mental_health_data)

# build_single_dataset(city_name, city_data, all_cities_df)

# print("\nMSE for different degree of Polynomials for Ground Truth:\n")
# best_degree_for_ground_truth =
↳fetch_best_degree_regression_model(city_data, normalize_original_data,
↳for_our_model=False)

# print("\nMSE for different degree of Polynomials for Our Model:\n")
# best_degree_for_our_model = fetch_best_degree_regression_model(city_data,
↳normalize_original_data)
best_degree_for_ground_truth, best_degree_for_our_model = 0, 0

print("\nBefore Covid (Mar-2020)\n\n")
city_data_subset = fetch_city_data_before_covid(city_data)
regress_ground_truth(city_name, city_data_subset,
↳best_degree_for_ground_truth, normalize_original_data, poly_model,
↳before_covid=True)
regress_our_model(city_name, city_data_subset, best_degree_for_our_model,
↳normalize_original_data, include_mental_health_data, poly_model,
↳before_covid=True)

```

```

    regress_our_model_without_ground_truth(city_name, city_data,
↪normalize_original_data, include_mental_health_data, poly_model,
↪before_covid=True, save_fig=True, time_period="WithoutGroundTruth")

    print("\nTill Dec-2021\n\n")
    regress_ground_truth(city_name, city_data, best_degree_for_ground_truth,
↪normalize_original_data, poly_model)
    regress_our_model(city_name, city_data, best_degree_for_our_model,
↪normalize_original_data, include_mental_health_data, poly_model)
    regress_our_model_without_ground_truth(city_name, city_data,
↪normalize_original_data, include_mental_health_data, poly_model,
↪before_covid=True, save_fig=True, time_period="WithoutGroundTruth")

    if city_name in ["LA", "CHICAGO", "NYC"]:
        infer_covid_comparison(city_name, city_data, absolute_file_path,
↪best_degree_for_our_model, normalize_original_data,
↪include_mental_health_data, poly_model)

    plot_correlation_graphs(city_data, normalize_original_data)

    plot_time_series_graph(city_name, city_data, save_fig=True)

```

```

[22]: # Create Final Dataset including all cities and run regression

# def infer_final_dataset_all_cities(all_cities_df):

#     all_cities_final_dataset = pd.concat(all_cities_df)
#     all_cities_final_dataset['MonthYear'] = pd.
↪to_datetime(all_cities_final_dataset['MonthYear'], format="%b-%y")
#     all_cities_final_dataset['CrimeCount'] = all_cities_final_dataset['Cul.
↪Crime count']
#     all_cities_final_dataset = all_cities_final_dataset.loc[:, ['MonthYear',
↪'City', 'CrimeCount', 'Gender_ratio', 'UnemploymentRate',
↪'MedianIncomeRate', 'WorkforceCount', 'PercentNegativeUsers',
↪'FavorOfDemocrats']]
#     all_cities_final_dataset = all_cities_final_dataset.
↪sort_values(by=['MonthYear', 'City'])

#     all_cities_final_dataset.to_csv('final_dataset.csv', index=False,
↪header=True)

#     def ols_coef(x, formula):
#         return smf.ols(formula, data=x).fit().params

```

```

#     gamma = (all_cities_final_dataset.groupby('MonthYear').
↳ apply(ols_coef, 'CrimeCount ~ 1 + Gender_ratio + UnemploymentRate +
↳ MedianIncomeRate + WorkforceCount + PercentNegativeUsers +
↳ FavorOfDemocrats'))
#     def fm_summary(p):
#         s = p.describe().T
#         s['std_error'] = s['std']/np.sqrt(s['count'])
#         s['tstat'] = s['mean']/s['std_error']
#         return s[['mean', 'std_error', 'tstat']]
#     print(fm_summary(gamma))

#     all_cities_final_dataset = all_cities_final_dataset.set_index(['City',
↳ 'MonthYear'])
#     print(all_cities_final_dataset)

#     model = FamaMacBeth.from_formula('CrimeCount ~ 1 + Gender_ratio +
↳ UnemploymentRate + MedianIncomeRate + WorkforceCount + PercentNegativeUsers
↳ + FavorOfDemocrats', data=all_cities_final_dataset)
#     result = model.fit(cov_type= 'kernel',debiased = False, bandwidth = 3)

#     print(result.summary)

```

```

[23]: # Plot Side-by-side bar graph for MSE in both the models

def plot_rmse(mse_our_model, std_our_model, mse_ground_truth, std_ground_truth,
↳ before_covid=False):
    mse_our_model_values, std_our_model_values, mse_ground_truth_values,
↳ std_ground_truth_values = [], [], [], []
    for city in cities:
        mse_our_model_values.append(mse_our_model[city])
        std_our_model_values.append(std_our_model[city])
        mse_ground_truth_values.append(mse_ground_truth[city])
        std_ground_truth_values.append(std_ground_truth[city])

    x_indices = np.arange(len(cities)) # the x locations for the groups
    width = 0.35 # the width of the bars

    fig = plt.figure()
    ax = fig.add_subplot(111)
    rects1 = ax.bar(x_indices, mse_ground_truth_values, width,
↳ color='lightblue', yerr=std_ground_truth_values)
    rects2 = ax.bar(x_indices+width, mse_our_model_values, width,
↳ color='royalblue', yerr=std_our_model_values)

    # add some
    ax.set_ylabel('Mean Square Error')

```



```

ax.set_title('MSE comparison of both the models')
ax.set_xlabel('Cities')
ax.set_xticks(x_indices + width / 2)
ax.set_xticklabels(cities, rotation=90)
ax.legend( (rects1[0], rects2[0]), ('Reference Paper', 'Our Model') )

#     if before_covid:
#         plt.savefig('RMSE_before_covid.png', transparent=True,
#             ↪bbox_inches='tight')
#     else:
#         plt.savefig('RMSE.png', transparent=True, bbox_inches='tight')

plt.show()

```

[24]: # Main Initiating Code Block

```

# cities_files = {'LA': 'CSV_LA_city_final_dataset_v2.csv', 'CHICAGO':
#     ↪'chicago_city_final_dataset_v2.csv', 'NYC': 'NYC_final_dataset_v2.csv'}
cities = ["NYC", "LA", "CHICAGO", "NASHVILLE", "BOSTON", "SEATTLE", "DENVER",
    ↪"DALLAS", "NEW_ORLEANS", "INDIANAPOLIS"]
# cities = ["LA"]
all_cities_df = []

for city_name in cities:
    print("\n\n{}:\n".format(city_name))
    city_file = "{}_final_dataset2.csv".format(city_name)
    infer_data(city_name, city_file, all_cities_df, absolute_file_path,
    ↪normalize_original_data=True, poly_model=False)

print("\nBefore Covid (Mar-2020)\n\n")
plot_rmse(mse_our_model_before_covid, std_our_model_before_covid,
    ↪mse_ground_truth_before_covid, std_ground_truth_before_covid,
    ↪before_covid=True)
print("\nTill Dec-2021\n\n")
plot_rmse(mse_our_model, std_our_model, mse_ground_truth, std_ground_truth)

# infer_final_dataset_all_cities(all_cities_df)

```

NYC:

Before Covid (Mar-2020)

Ground Truth:

Linear Model summary:

OLS Regression Results

```

=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.834
Model:              OLS                    Adj. R-squared:  0.829
Method:             Least Squares          F-statistic:    177.7
Date:               Sun, 11 Dec 2022        Prob (F-statistic): 3.32e-41
Time:               16:45:20                Log-Likelihood:  119.19
No. Observations:   110                    AIC:            -230.4
Df Residuals:       106                    BIC:            -219.6
Df Model:           3
Covariance Type:    nonrobust
=====

```

```

=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
const                27.2471      16.375        1.664      0.099      -5.218
59.712
Gender_ratio         -50.7675      31.349       -1.619      0.108     -112.919
11.384
Scaled_Unemployment    0.5384       0.173        3.109      0.002        0.195
0.882
Scaled_MedianIncomeRate -0.3993      0.061       -6.541      0.000       -0.520
-0.278
=====
Omnibus:              35.741      Durbin-Watson:      1.219
Prob(Omnibus):         0.000      Jarque-Bera (JB):    90.769
Skew:                  -1.198      Prob(JB):            1.95e-20
Kurtosis:               6.750      Cond. No.            5.39e+03
=====


```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.39e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:          Scaled_CrimeCount      R-squared:                0.842
Model:                  OLS                   Adj. R-squared:           0.831
Method:                 Least Squares         F-statistic:             77.80
Date:                  Sun, 11 Dec 2022       Prob (F-statistic):      4.98e-38
Time:                  16:45:22              Log-Likelihood:          121.94
No. Observations:      110                  AIC:                    -227.9
Df Residuals:          102                  BIC:                    -206.3
Df Model:              7
Covariance Type:       nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                17.2737    44.739      0.386    0.700
-71.466    106.013
Gender_ratio        -33.3607    88.219     -0.378    0.706
-208.343    141.621
Scaled_Unemployment    0.5347     0.217      2.465    0.015
0.104     0.965
Scaled_MedianIncomeRate -0.3891     0.063     -6.224    0.000
-0.513    -0.265
WorkforceCount      -7.6674    14.337     -0.535    0.594
=====

```


-36.105	20.770				
PercentNegativeUsers		-0.1104	0.135	-0.817	0.416
-0.378	0.158				
FavorOfDemocrats		1.5674	2.028	0.773	0.441
-2.454	5.589				
Scaled_SearchCountForDepression		0.0834	0.047	1.778	0.078
-0.010	0.176				
=====					
Omnibus:	32.899	Durbin-Watson:		1.254	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		75.456	
Skew:	-1.142	Prob(JB):		4.12e-17	
Kurtosis:	6.353	Cond. No.		1.94e+04	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.94e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

OLS Regression Results

Dep. Variable:	Scaled_CrimeCount	R-squared:	0.438
----------------	-------------------	------------	-------

```

Model:                OLS      Adj. R-squared:      0.420
Method:               Least Squares    F-statistic:      24.72
Date:                Sun, 11 Dec 2022    Prob (F-statistic): 3.80e-15
Time:                16:45:22    Log-Likelihood:    29.632
No. Observations:    132    AIC:                -49.26
Df Residuals:        127    BIC:                -34.85
Df Model:             4
Covariance Type:     nonrobust

```

```

=====
=====

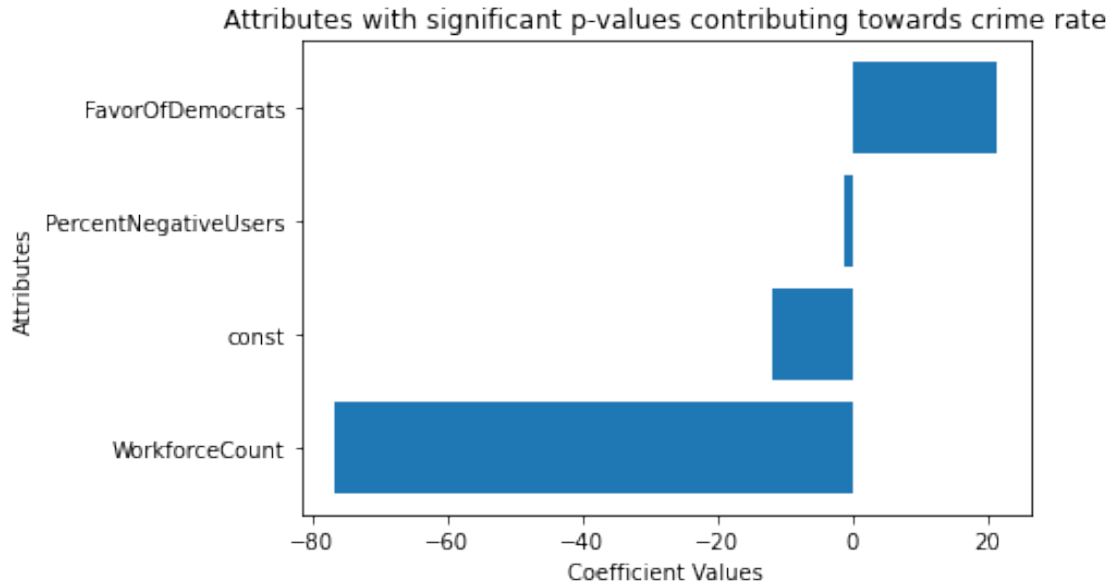
```

	coef	std err	t	P> t
[0.025      0.975]				
-----				
const	-11.9249	2.214	-5.386	0.000
-16.306      -7.543				
WorkforceCount	-76.7194	9.522	-8.057	0.000
-95.561      -57.878				
PercentNegativeUsers	-1.2691	0.296	-4.293	0.000
-1.854      -0.684				
FavorOfDemocrats	21.5093	3.085	6.972	0.000
15.404      27.614				
Scaled_SearchCountForDepression	-0.1113	0.089	-1.251	0.213
-0.287      0.065				
=====				
Omnibus:	18.806	Durbin-Watson:	0.460	
Prob(Omnibus):	0.000	Jarque-Bera (JB):	25.765	
Skew:	0.770	Prob(JB):	2.54e-06	
Kurtosis:	4.521	Cond. No.	764.	
=====				

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Till Dec-2021

Ground Truth:

Linear Model summary:

OLS Regression Results					
=====					
Dep. Variable:	Scaled_CrimeCount	R-squared:	0.885		
Model:	OLS	Adj. R-squared:	0.882		
Method:	Least Squares	F-statistic:	327.8		
Date:	Sun, 11 Dec 2022	Prob (F-statistic):	7.19e-60		
Time:	16:45:22	Log-Likelihood:	134.28		
No. Observations:	132	AIC:	-260.6		
Df Residuals:	128	BIC:	-249.0		
Df Model:	3				
Covariance Type:	nonrobust				
=====					
=====					
	coef	std err	t	P> t	[0.025
0.975]					
-----					
const	-9.7313	5.039	-1.931	0.056	-19.702
0.240					
Gender_ratio	20.3820	9.620	2.119	0.036	1.347
39.417					
Scaled_Unemployment	-0.2677	0.043	-6.162	0.000	-0.354

-0.182  
 Scaled\_MedianIncomeRate -0.6854 0.028 -24.075 0.000 -0.742  
 -0.629

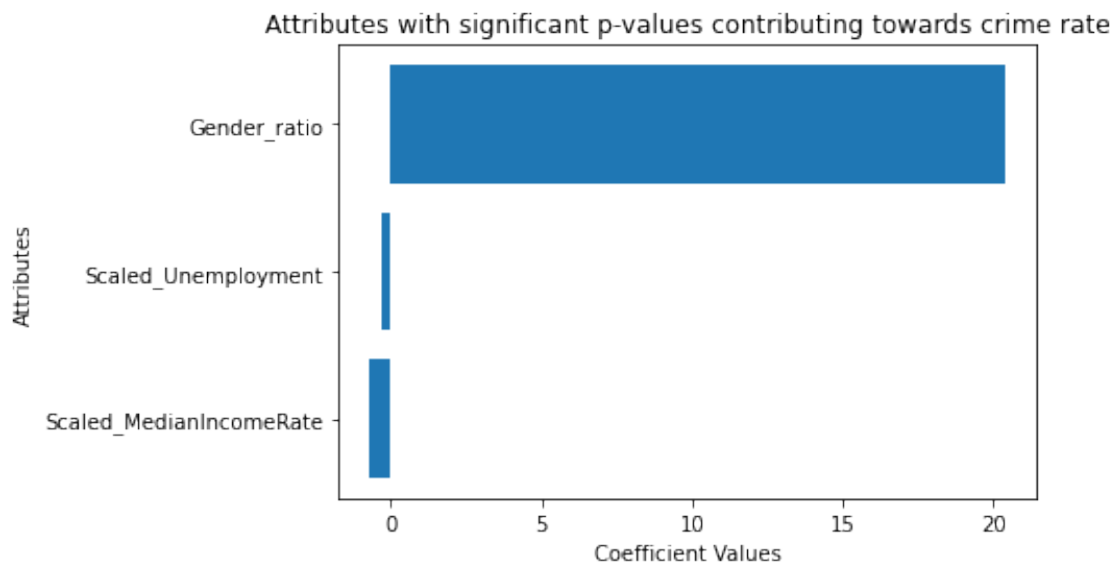
```
=====
Omnibus:                31.872    Durbin-Watson:                1.179
Prob(Omnibus):           0.000    Jarque-Bera (JB):           74.247
Skew:                   -0.964    Prob(JB):                   7.54e-17
Kurtosis:                6.128    Cond. No.                   1.76e+03
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.76e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

OLS Regression Results

```
=====
Dep. Variable:    Scaled_CrimeCount    R-squared:                0.890
Model:            OLS                 Adj. R-squared:            0.884
Method:           Least Squares        F-statistic:              143.6
Date:             Sun, 11 Dec 2022     Prob (F-statistic):       2.41e-56
Time:             16:45:22             Log-Likelihood:           137.43
=====
```

```

No. Observations:      132    AIC:                -258.9
Df Residuals:          124    BIC:                -235.8
Df Model:               7
Covariance Type:       nonrobust

```

```

=====

```

		coef	std err	t	P> t
-----					
	[0.025      0.975]				
-----					
const		-5.9746	5.525	-1.081	0.282
	-16.911      4.961				
Gender_ratio		11.5131	10.661	1.080	0.282
	-9.589      32.615				
Scaled_Unemployment		-0.2438	0.048	-5.101	0.000
	-0.338      -0.149				
Scaled_MedianIncomeRate		-0.6702	0.033	-20.346	0.000
	-0.735      -0.605				
WorkforceCount		-7.3509	5.573	-1.319	0.190
	-18.381      3.679				
PercentNegativeUsers		-0.2464	0.140	-1.760	0.081
	-0.524      0.031				
FavorOfDemocrats		1.6608	1.709	0.972	0.333
	-1.722      5.044				
Scaled_SearchCountForDepression		0.0531	0.042	1.258	0.211
	-0.031      0.137				
=====					
Omnibus:	31.465	Durbin-Watson:	1.200		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	75.290		
Skew:	-0.940	Prob(JB):	4.48e-17		
Kurtosis:	6.187	Cond. No.	2.47e+03		
=====					


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.47e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:





Our Model without ground truth:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:          Scaled_CrimeCount      R-squared:                0.438
Model:                  OLS                   Adj. R-squared:           0.420
Method:                 Least Squares         F-statistic:              24.72
Date:                  Sun, 11 Dec 2022       Prob (F-statistic):       3.80e-15
Time:                  16:45:22               Log-Likelihood:           29.632
No. Observations:      132                   AIC:                     -49.26
Df Residuals:          127                   BIC:                     -34.85
Df Model:               4
Covariance Type:       nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                -11.9249         2.214     -5.386     0.000
-16.306    -7.543
WorkforceCount       -76.7194         9.522     -8.057     0.000
-95.561   -57.878
PercentNegativeUsers   -1.2691         0.296     -4.293     0.000
-1.854    -0.684
FavorOfDemocrats      21.5093         3.085      6.972     0.000
15.404    27.614
Scaled_SearchCountForDepression -0.1113         0.089     -1.251     0.213
=====

```


-0.287      0.065

```
=====
Omnibus:                18.806   Durbin-Watson:                0.460
Prob(Omnibus):          0.000   Jarque-Bera (JB):            25.765
Skew:                   0.770   Prob(JB):                    2.54e-06
Kurtosis:               4.521   Cond. No.                     764.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Covid Data Comparison:

Regression from Mar-2020 to Dec-2021 along with Covid Positive Cases Per Month:

Our Model:

Linear Model summary:

```

                        OLS Regression Results
=====
Dep. Variable:          Scaled_CrimeCount   R-squared:                0.568
Model:                  OLS                 Adj. R-squared:           0.432
Method:                 Least Squares        F-statistic:             4.200
```

```

Date:                Sun, 11 Dec 2022    Prob (F-statistic):        0.0125
Time:                16:45:23           Log-Likelihood:          39.654
No. Observations:    22                 AIC:                    -67.31
Df Residuals:        16                 BIC:                    -60.76
Df Model:            5
Covariance Type:     nonrobust

```

```

=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
Gender_ratio                1.0775      0.748      1.440      0.169
-0.508      2.663
Scaled_Unemployment        -0.1864      0.063     -2.954      0.009
-0.320     -0.053
Scaled_MedianIncomeRate    -0.7378      0.819     -0.901      0.381
-2.475      0.999
WorkforceCount             -0.0758      0.071     -1.067      0.302
-0.226      0.075
PercentNegativeUsers       -1.5430      0.834     -1.850      0.083
-3.311      0.225
FavorOfDemocrats           1.3225      0.895      1.478      0.159
-0.575      3.220
Scaled_SearchCountForDepression  0.1518      0.084      1.802      0.090
-0.027      0.330
CovidPercPositive          -0.2757      0.148     -1.859      0.082
-0.590      0.039
=====
Omnibus:                    0.969    Durbin-Watson:          2.645
Prob(Omnibus):              0.616    Jarque-Bera (JB):        0.409
Skew:                      0.334    Prob(JB):                0.815
Kurtosis:                  3.027    Cond. No.                5.73e+17
=====

```

#### Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 1.41e-34. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Plot for Linear Model Coefficients for Our Model:


```

/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1807615713.py:1
9: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.

```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
city_combined_dataset_covid['CovidPercPositive'] = perc
```



Regression for the same time frame without considering Covid Cases:

Our Model:

Linear Model summary:

### OLS Regression Results

Dep. Variable:	Scaled_CrimeCount	R-squared:	0.474
Model:	OLS	Adj. R-squared:	0.350
Method:	Least Squares	F-statistic:	3.833
Date:	Sun, 11 Dec 2022	Prob (F-statistic):	0.0213
Time:	16:45:23	Log-Likelihood:	37.503
No. Observations:	22	AIC:	-65.01
Df Residuals:	17	BIC:	-59.55
Df Model:	4		
Covariance Type:	nonrobust		

		coef	std err	t	P> t
[0.025	0.975]				

Gender_ratio	0.5494	0.740	0.742	0.468
-1.012	2.111			
Scaled_Unemployment	-0.1175	0.055	-2.151	0.046
-0.233	-0.002			
Scaled_MedianIncomeRate	-0.1462	0.808	-0.181	0.858
-1.850	1.558			
WorkforceCount	-0.0246	0.070	-0.351	0.730
-0.172	0.123			
PercentNegativeUsers	-1.1112	0.857	-1.297	0.212
-2.919	0.696			
FavorOfDemocrats	0.6929	0.886	0.782	0.445
-1.177	2.563			
Scaled_SearchCountForDepression	0.0247	0.053	0.469	0.645
-0.086	0.136			

---

Omnibus:	2.053	Durbin-Watson:	2.341
Prob(Omnibus):	0.358	Jarque-Bera (JB):	0.899
Skew:	0.465	Prob(JB):	0.638
Kurtosis:	3.342	Cond. No.	2.97e+18


---

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 5.23e-36. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Plot for Linear Model Coefficients for Our Model:



## Correlation Graphs:




```
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
1: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "g-" (-> color='g'). The keyword argument will take precedence.
```

```
ax1.plot(city_data["MonthYear"], city_data['Scaled_Unemployment'], 'g-',
color="#6495ED")
```

```
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
2: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "b-" (-> color='b'). The keyword argument will take precedence.
```

```
ax2.plot(city_data["MonthYear"], city_data['Scaled_CrimeCount'], 'b-',
color="#CC4F1B")
```

```
<Figure size 2400x1200 with 0 Axes>
```



LA:

Before Covid (Mar-2020)

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.864
Model:              OLS                    Adj. R-squared:  0.860
Method:             Least Squares          F-statistic:    224.9
Date:               Sun, 11 Dec 2022        Prob (F-statistic): 8.48e-46
Time:               16:45:23                Log-Likelihood: 119.41
No. Observations:   110                    AIC:            -230.8
Df Residuals:       106                    BIC:            -220.0
Df Model:           3
Covariance Type:    nonrobust
=====
```

```
=====
coef      std err      t      P>|t|      [0.025
```


0.975]

-----					
-----					
const	-0.2077	2.189	-0.095	0.925	-4.547
4.132					
Gender_ratio	1.3841	4.296	0.322	0.748	-7.134
9.902					
Scaled_Unemployment	0.7707	0.108	7.155	0.000	0.557
0.984					
Scaled_MedianIncomeRate	-0.2302	0.065	-3.562	0.001	-0.358
-0.102					
=====					
Omnibus:	0.714	Durbin-Watson:		1.177	
Prob(Omnibus):	0.700	Jarque-Bera (JB):		0.312	
Skew:	0.049	Prob(JB):		0.856	
Kurtosis:	3.242	Cond. No.		725.	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

OLS Regression Results



```

=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.890
Model:              OLS                    Adj. R-squared:  0.883
Method:              Least Squares         F-statistic:     118.5
Date:                Sun, 11 Dec 2022      Prob (F-statistic): 4.75e-46
Time:                16:45:23              Log-Likelihood:  131.23
No. Observations:    110                   AIC:             -246.5
Df Residuals:        102                   BIC:             -224.8
Df Model:            7
Covariance Type:     nonrobust
=====

=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
const                -6.4065         2.537        -2.525     0.013
-11.438      -1.375
Gender_ratio          9.9219         4.337         2.287     0.024
1.319       18.525
Scaled_Unemployment   1.2847         0.145         8.869     0.000
0.997       1.572
Scaled_MedianIncomeRate -0.1025         0.073        -1.409     0.162
-0.247       0.042
WorkforceCount       -30.6813         6.485        -4.731     0.000
-43.544      -17.818
PercentNegativeUsers   0.0555         0.063         0.880     0.381
-0.070       0.181
FavorOfDemocrats      3.3964         1.121         3.029     0.003
1.173       5.620
Scaled_SearchCountForDepression 0.0823         0.045         1.817     0.072
-0.008       0.172
=====

Omnibus:              3.685    Durbin-Watson:      1.202
Prob(Omnibus):        0.158    Jarque-Bera (JB):    3.051
Skew:                 -0.349    Prob(JB):            0.217
Kurtosis:             3.422    Cond. No.            1.48e+03
=====

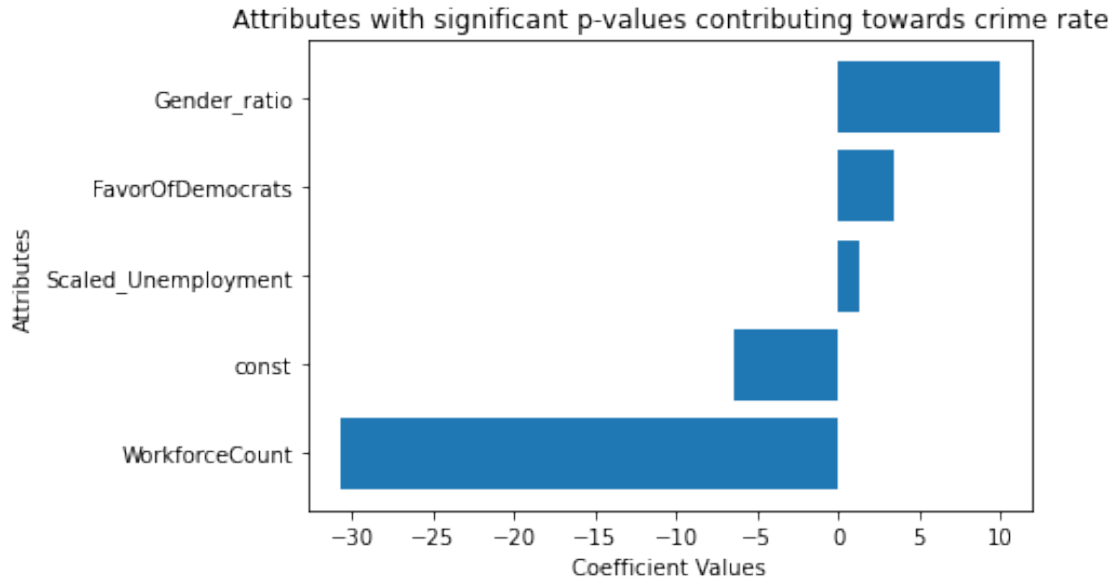
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.48e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:          Scaled_CrimeCount      R-squared:                0.575
Model:                  OLS                   Adj. R-squared:           0.562
Method:                 Least Squares          F-statistic:              43.04
Date:                   Sun, 11 Dec 2022        Prob (F-statistic):       8.83e-23
Time:                   16:45:24               Log-Likelihood:           41.755
No. Observations:       132                   AIC:                     -73.51
Df Residuals:           127                   BIC:                     -59.10
Df Model:                4
Covariance Type:        nonrobust
=====

```

```

=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
const                                8.0421      0.954      8.429      0.000
6.154      9.930
WorkforceCount                     -30.9226      9.846     -3.140      0.002
-50.407     -11.438
PercentNegativeUsers                 -0.2529      0.145     -1.750      0.083
-0.539      0.033
FavorOfDemocrats                    -9.4381      1.518     -6.216      0.000
-12.443     -6.433
Scaled_SearchCountForDepression     -0.1676      0.086     -1.939      0.055
=====

```


-0.339      0.003

```
=====
Omnibus:                5.613    Durbin-Watson:                0.417
Prob(Omnibus):          0.060    Jarque-Bera (JB):            2.849
Skew:                  -0.048    Prob(JB):                    0.241
Kurtosis:              2.287    Cond. No.                     873.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Till Dec-2021

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:    Scaled_CrimeCount    R-squared:                0.857
Model:            OLS                  Adj. R-squared:            0.854
Method:           Least Squares        F-statistic:              256.3
=====
```

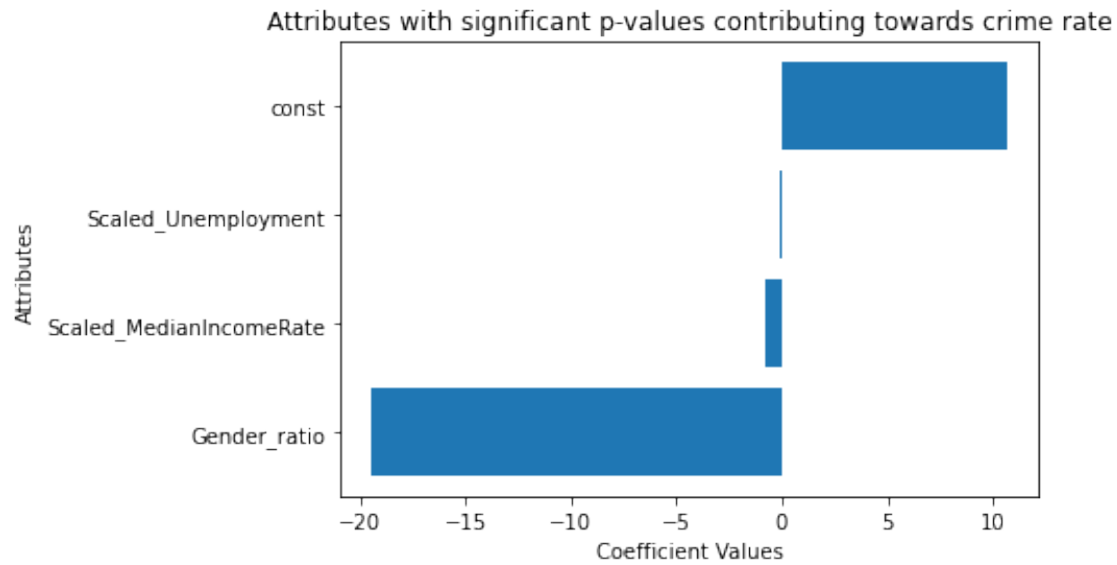
Date: Sun, 11 Dec 2022 Prob (F-statistic): 6.51e-54  
 Time: 16:45:24 Log-Likelihood: 113.70  
 No. Observations: 132 AIC: -219.4  
 Df Residuals: 128 BIC: -207.9  
 Df Model: 3  
 Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025
0.975]					
-----					
-----					
const	10.6498	2.114	5.037	0.000	6.466
14.834					
Gender_ratio	-19.4411	4.189	-4.641	0.000	-27.730
-11.153					
Scaled_Unemployment	-0.1041	0.044	-2.376	0.019	-0.191
-0.017					
Scaled_MedianIncomeRate	-0.7496	0.028	-26.773	0.000	-0.805
-0.694					
=====					
Omnibus:	0.100	Durbin-Watson:		0.969	
Prob(Omnibus):	0.951	Jarque-Bera (JB):		0.239	
Skew:	0.046	Prob(JB):		0.887	
Kurtosis:	2.813	Cond. No.		642.	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:    Scaled_CrimeCount    R-squared:            0.858
Model:            OLS                  Adj. R-squared:       0.850
Method:           Least Squares        F-statistic:          107.1
Date:             Sun, 11 Dec 2022     Prob (F-statistic):   1.78e-49
Time:             16:45:24             Log-Likelihood:       114.08
No. Observations: 132                  AIC:                  -212.2
Df Residuals:     124                  BIC:                  -189.1
Df Model:         7
Covariance Type:  nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                10.9824    2.320        4.733    0.000
6.390    15.575
Gender_ratio        -19.1958    4.336       -4.428    0.000
-27.777   -10.614
Scaled_Unemployment  -0.1280    0.056       -2.282    0.024
-0.239    -0.017
Scaled_MedianIncomeRate -0.7417    0.048     -15.383    0.000
-0.837    -0.646
WorkforceCount       4.4923    6.928        0.648    0.518
```


-9.221	18.205				
PercentNegativeUsers		-0.0257	0.086	-0.299	0.765
-0.195	0.144				
FavorOfDemocrats		-0.7881	1.234	-0.638	0.524
-3.231	1.655				
Scaled_SearchCountForDepression		0.0128	0.053	0.242	0.809
-0.092	0.117				
=====					
Omnibus:	0.103	Durbin-Watson:		0.980	
Prob(Omnibus):	0.950	Jarque-Bera (JB):		0.173	
Skew:	0.065	Prob(JB):		0.917	
Kurtosis:	2.879	Cond. No.		1.19e+03	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.19e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

OLS Regression Results

Dep. Variable:	Scaled_CrimeCount	R-squared:	0.575
----------------	-------------------	------------	-------

```

Model:                OLS      Adj. R-squared:      0.562
Method:               Least Squares    F-statistic:      43.04
Date:                Sun, 11 Dec 2022    Prob (F-statistic): 8.83e-23
Time:                16:45:24    Log-Likelihood:    41.755
No. Observations:    132    AIC:                -73.51
Df Residuals:        127    BIC:                -59.10
Df Model:             4
Covariance Type:     nonrobust

```

```

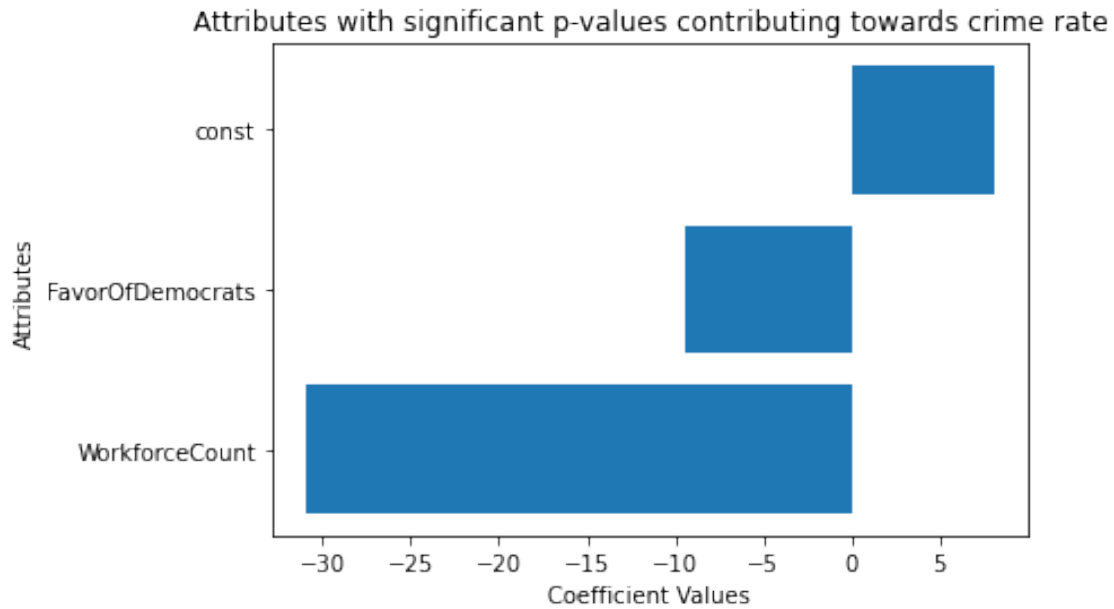
=====
=====
                                coef      std err          t      P>|t|
[0.025      0.975]
-----
const                8.0421      0.954      8.429      0.000
6.154      9.930
WorkforceCount       -30.9226      9.846     -3.140      0.002
-50.407     -11.438
PercentNegativeUsers  -0.2529      0.145     -1.750      0.083
-0.539      0.033
FavorOfDemocrats     -9.4381      1.518     -6.216      0.000
-12.443     -6.433
Scaled_SearchCountForDepression -0.1676      0.086     -1.939      0.055
-0.339      0.003
=====
Omnibus:              5.613    Durbin-Watson:      0.417
Prob(Omnibus):        0.060    Jarque-Bera (JB):    2.849
Skew:                 -0.048    Prob(JB):            0.241
Kurtosis:             2.287    Cond. No.            873.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Covid Data Comparison:

Regression from Mar-2020 to Dec-2021 along with Covid Positive Cases Per Month:

Our Model:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:    Scaled_CrimeCount    R-squared:        0.256
Model:            OLS                  Adj. R-squared:    0.024
Method:           Least Squares        F-statistic:       1.101
Date:             Sun, 11 Dec 2022     Prob (F-statistic): 0.398
Time:             16:45:24             Log-Likelihood:    33.988
No. Observations: 22                  AIC:               -55.98
Df Residuals:     16                  BIC:               -49.43
Df Model:         5
Covariance Type:  nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
Gender_ratio                0.0420    0.291      0.145      0.887
-0.574    0.658
Scaled_Unemployment         0.0262    0.066      0.395      0.698

```



-0.114	0.167				
Scaled_MedianIncomeRate		0.1545	0.211	0.731	0.475
-0.293	0.602				
WorkforceCount		0.0026	0.011	0.234	0.818
-0.021	0.026				
PercentNegativeUsers		-0.2587	0.639	-0.405	0.691
-1.612	1.095				
FavorOfDemocrats		0.0614	0.393	0.156	0.878
-0.771	0.894				
Scaled_SearchCountForDepression		0.0706	0.086	0.817	0.426
-0.113	0.254				
CovidPercPositive		-0.4031	0.351	-1.147	0.268
-1.148	0.342				

Omnibus:	0.621	Durbin-Watson:	1.761
Prob(Omnibus):	0.733	Jarque-Bera (JB):	0.638
Skew:	-0.093	Prob(JB):	0.727
Kurtosis:	2.187	Cond. No.	2.90e+17

#### Notes:


- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 6.03e-34. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

#### Plot for Linear Model Coefficients for Our Model:

```
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1807615713.py:1
9: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
city_combined_dataset_covid['CovidPercPositive'] = perc
```



Regression for the same time frame without considering Covid Cases:

Our Model:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:          Scaled_CrimeCount    R-squared:                 0.195
Model:                  OLS                 Adj. R-squared:            0.005
Method:                 Least Squares        F-statistic:               1.028
Date:                  Sun, 11 Dec 2022      Prob (F-statistic):       0.421
Time:                  16:45:24             Log-Likelihood:           33.118
No. Observations:      22                  AIC:                      -56.24
Df Residuals:          17                  BIC:                      -50.78
Df Model:              4
Covariance Type:       nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
Gender_ratio                -0.0043    0.291      -0.015    0.988
-0.617    0.609
Scaled_Unemployment          0.0309    0.067       0.462    0.650
```

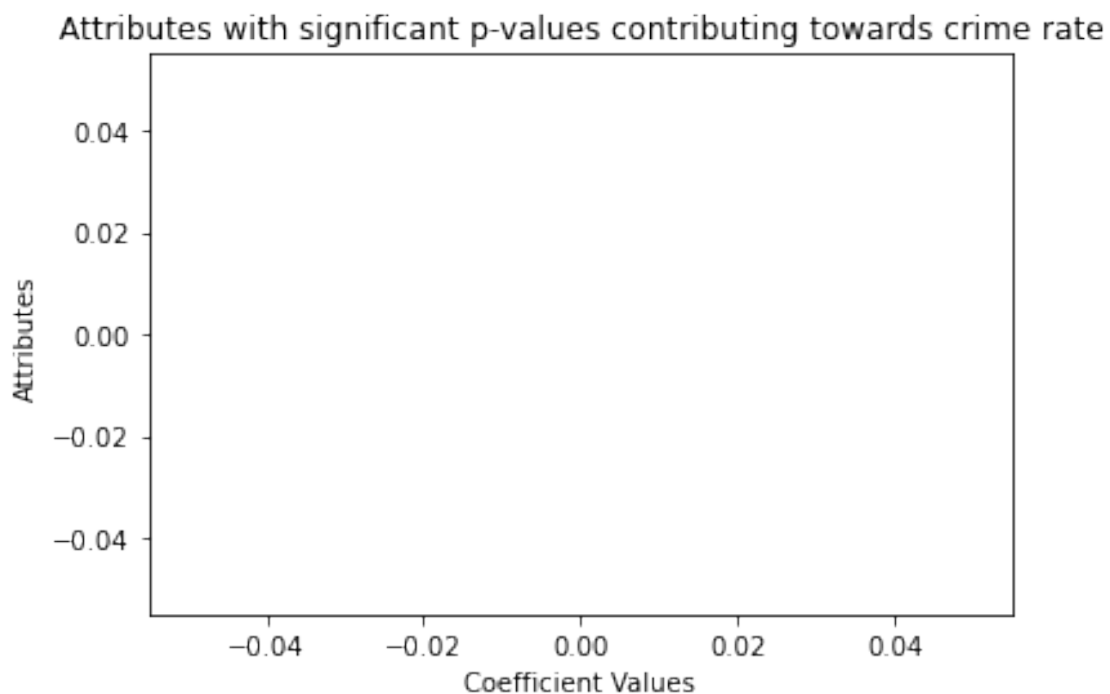
-0.110	0.172				
Scaled_MedianIncomeRate		0.2549	0.194	1.314	0.206
-0.154	0.664				
WorkforceCount		0.0012	0.011	0.112	0.912
-0.022	0.024				
PercentNegativeUsers		-0.3802	0.636	-0.598	0.558
-1.721	0.961				
FavorOfDemocrats		0.0008	0.393	0.002	0.998
-0.828	0.829				
Scaled_SearchCountForDepression		0.0879	0.086	1.023	0.320
-0.093	0.269				
=====					
Omnibus:	2.374	Durbin-Watson:	1.756		
Prob(Omnibus):	0.305	Jarque-Bera (JB):	1.193		
Skew:	-0.151	Prob(JB):	0.551		
Kurtosis:	1.900	Cond. No.	3.81e+17		
=====					

Notes:

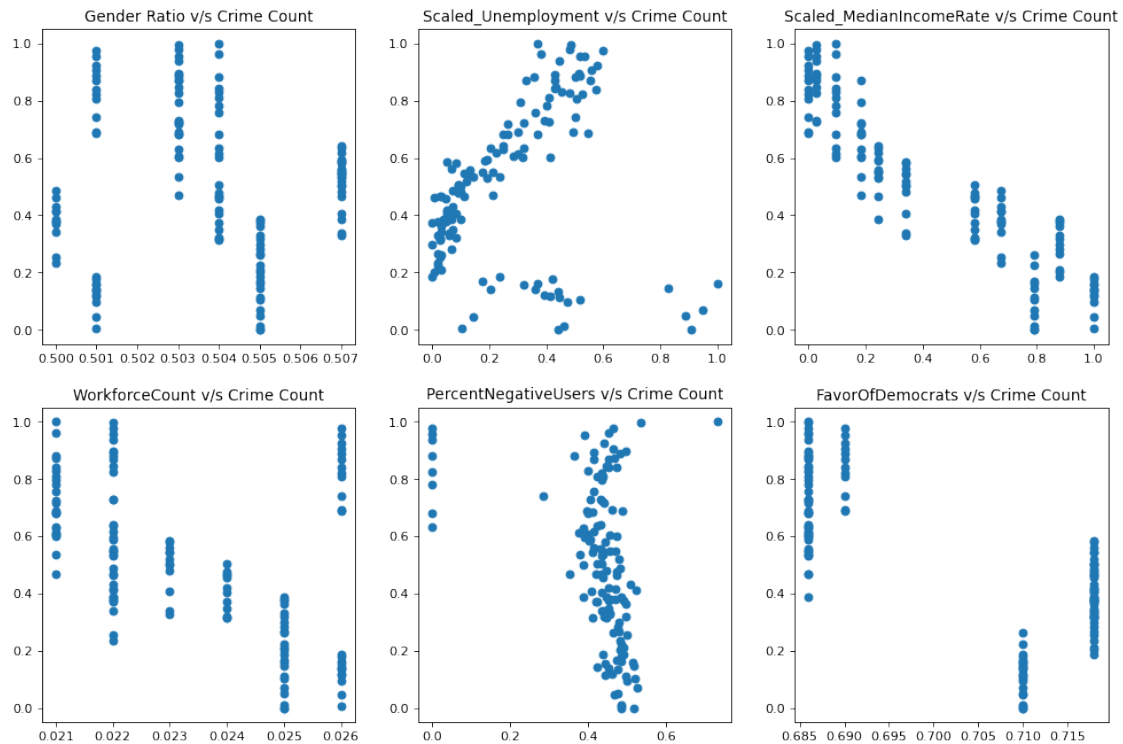
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 3.48e-34. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Plot for Linear Model Coefficients for Our Model:



## Correlation Graphs:



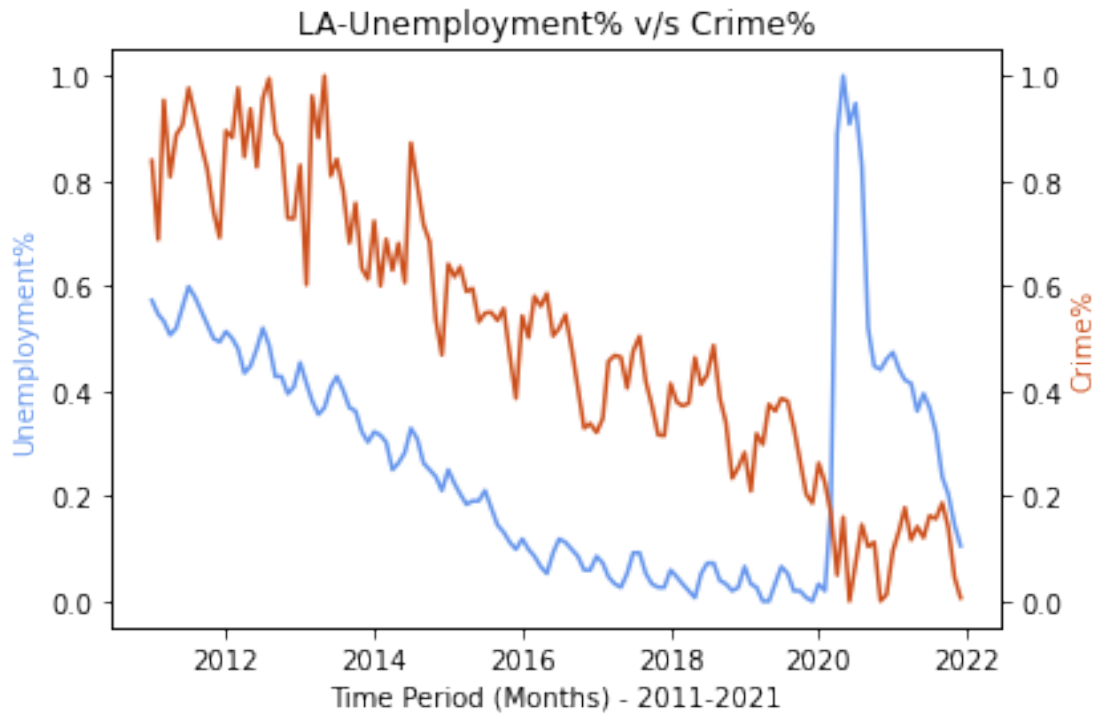
```
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
1: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "g-" (-> color='g'). The keyword argument will take precedence.
```

```
ax1.plot(city_data["MonthYear"], city_data['Scaled_Unemployment'], 'g-',
color="#6495ED")
```

```
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
2: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "b-" (-> color='b'). The keyword argument will take precedence.
```

```
ax2.plot(city_data["MonthYear"], city_data['Scaled_CrimeCount'], 'b-',
color="#CC4F1B")
```

```
<Figure size 2400x1200 with 0 Axes>
```



CHICAGO:

Before Covid (Mar-2020)

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.777
Model:              OLS                    Adj. R-squared:  0.770
Method:             Least Squares          F-statistic:    122.8
Date:               Sun, 11 Dec 2022        Prob (F-statistic): 2.33e-34
Time:               16:45:25                Log-Likelihood:  105.41
No. Observations:   110                    AIC:            -202.8
Df Residuals:       106                    BIC:            -192.0
Df Model:           3
Covariance Type:    nonrobust
=====
```

```
=====
coef      std err      t      P>|t|      [0.025
```


0.975]

const	3.8112	3.316	1.149	0.253	-2.764
10.386					
Gender_ratio	-6.5741	6.382	-1.030	0.305	-19.227
6.079					
Scaled_Unemployment	0.8608	0.173	4.986	0.000	0.519
1.203					
Scaled_MedianIncomeRate	-0.1702	0.108	-1.573	0.119	-0.385
0.044					
=====					
Omnibus:	5.500	Durbin-Watson:		0.573	
Prob(Omnibus):	0.064	Jarque-Bera (JB):		3.443	
Skew:	-0.252	Prob(JB):		0.179	
Kurtosis:	2.295	Cond. No.		960.	

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

# OLS Regression Results

```

=====
Dep. Variable:          Scaled_CrimeCount      R-squared:                0.805
Model:                  OLS                    Adj. R-squared:           0.791
Method:                 Least Squares          F-statistic:              60.05
Date:                   Sun, 11 Dec 2022        Prob (F-statistic):       2.37e-33
Time:                   16:45:25               Log-Likelihood:           112.82
No. Observations:       110                   AIC:                     -209.6
Df Residuals:           102                   BIC:                     -188.0
Df Model:                7
Covariance Type:        nonrobust
=====

```

```

=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
const                        1.8651      4.017        0.464      0.643
-6.103      9.833
Gender_ratio                 -8.7363      6.720       -1.300      0.197
-22.065      4.593
Scaled_Unemployment          0.6259      0.177        3.532      0.001
0.274      0.977
Scaled_MedianIncomeRate     -0.2615      0.122       -2.149      0.034
-0.503     -0.020
WorkforceCount             -15.2992      7.122       -2.148      0.034
-29.425     -1.173
PercentNegativeUsers         0.1180      0.107        1.104      0.272
-0.094      0.330
FavorOfDemocrats            5.0899      2.230        2.283      0.025
0.667      9.513
Scaled_SearchCountForDepression -0.1202      0.049       -2.477      0.015
-0.216     -0.024
=====
Omnibus:                    4.584      Durbin-Watson:            0.612
Prob(Omnibus):              0.101      Jarque-Bera (JB):         2.888
Skew:                       -0.205      Prob(JB):                 0.236
Kurtosis:                   2.320      Cond. No.                  1.47e+03
=====

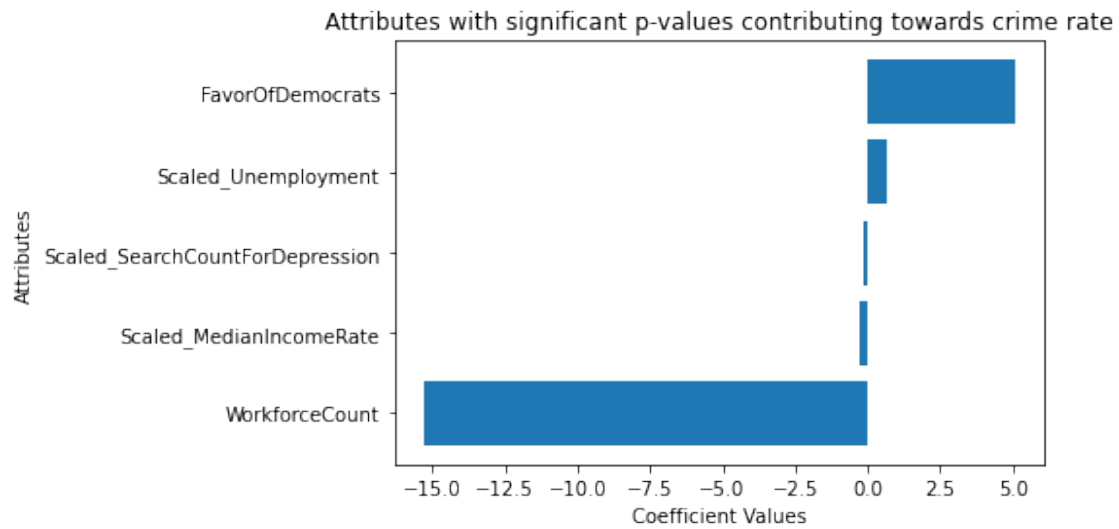
```

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.47e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:    Scaled_CrimeCount    R-squared:                0.235
Model:            OLS                  Adj. R-squared:           0.211
Method:           Least Squares        F-statistic:              9.769
Date:             Sun, 11 Dec 2022      Prob (F-statistic):       6.37e-07
Time:             16:45:25              Log-Likelihood:           12.958
No. Observations: 132                  AIC:                     -15.92
Df Residuals:     127                  BIC:                     -1.503
Df Model:          4
Covariance Type:  nonrobust
=====
```

		coef	std err	t	P> t
-----					
	[0.025      0.975]				
-----					
const		-11.6414	3.309	-3.518	0.001
	-18.189      -5.094				
WorkforceCount		-18.9933	14.044	-1.352	0.179
	-46.785      8.798				
PercentNegativeUsers		0.3077	0.251	1.228	0.222
	-0.188      0.803				
FavorOfDemocrats		17.4337	5.001	3.486	0.001
	7.538      27.329				
Scaled_SearchCountForDepression		-0.3310	0.092	-3.608	0.000



-0.513      -0.149

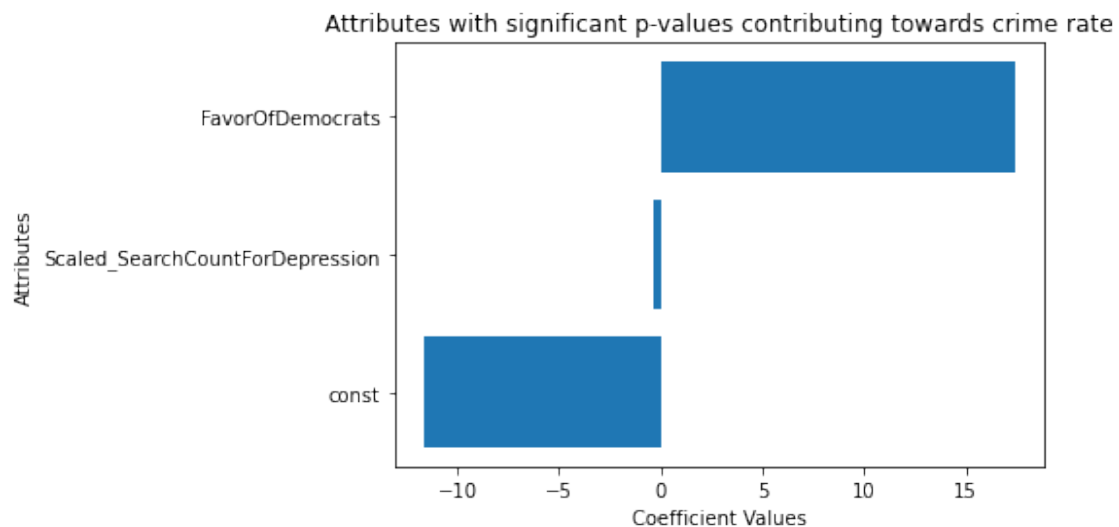
```
=====
Omnibus:                5.838    Durbin-Watson:                0.144
Prob(Omnibus):          0.054    Jarque-Bera (JB):          6.010
Skew:                   -0.510    Prob(JB):                  0.0495
Kurtosis:               2.774    Cond. No.                  1.05e+03
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.05e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model without ground truth:



Till Dec-2021

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:          Scaled_CrimeCount    R-squared:                0.842
Model:                  OLS                  Adj. R-squared:           0.839
Method:                 Least Squares        F-statistic:              227.8
Date:                   Sun, 11 Dec 2022     Prob (F-statistic):       3.87e-51
=====
```

```

Time:                  16:45:25   Log-Likelihood:          117.15
No. Observations:      132       AIC:                  -226.3
Df Residuals:          128       BIC:                  -214.8
Df Model:               3
Covariance Type:       nonrobust

```

```


=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const              12.1945      3.184      3.830      0.000      5.894
18.495
Gender_ratio       -22.0308      6.187     -3.561      0.001     -34.273
-9.788
Scaled_Unemployment -0.1000      0.057     -1.766      0.080     -0.212
0.012
Scaled_MedianIncomeRate -0.7663      0.032    -24.223      0.000     -0.829
-0.704
=====
Omnibus:           0.946   Durbin-Watson:           0.554
Prob(Omnibus):     0.623   Jarque-Bera (JB):           0.959
Skew:              0.060   Prob(JB):                 0.619
Kurtosis:          2.600   Cond. No.                  980.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:          Scaled_CrimeCount      R-squared:                0.872
Model:                  OLS                   Adj. R-squared:           0.865
Method:                 Least Squares         F-statistic:             120.5
Date:                   Sun, 11 Dec 2022       Prob (F-statistic):      3.33e-52
Time:                   16:45:25              Log-Likelihood:          130.85
No. Observations:       132                   AIC:                    -245.7
Df Residuals:           124                   BIC:                    -222.6
Df Model:                7
Covariance Type:        nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                8.2253    3.782        2.175    0.032
0.740    15.711
Gender_ratio        -21.8722    5.990       -3.652    0.000
-33.727   -10.017
Scaled_Unemployment  -0.1198    0.058       -2.053    0.042
-0.235    -0.004
Scaled_MedianIncomeRate  -0.7440    0.030      -24.400    0.000
-0.804    -0.684
WorkforceCount      -23.4736    5.852       -4.011    0.000

```


-35.056	-11.891				
PercentNegativeUsers		0.2090	0.105	1.988	0.049
0.001	0.417				
FavorOfDemocrats		6.4892	2.224	2.918	0.004
2.087	10.891				
Scaled_SearchCountForDepression		-0.1277	0.041	-3.150	0.002
-0.208	-0.047				
=====					
Omnibus:	0.423	Durbin-Watson:		0.626	
Prob(Omnibus):	0.809	Jarque-Bera (JB):		0.572	
Skew:	-0.017	Prob(JB):		0.751	
Kurtosis:	2.679	Cond. No.		1.39e+03	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.39e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

OLS Regression Results

Dep. Variable:	Scaled_CrimeCount	R-squared:	0.235
Model:	OLS	Adj. R-squared:	0.211

```

Method:                Least Squares    F-statistic:                9.769
Date:                  Sun, 11 Dec 2022  Prob (F-statistic):        6.37e-07
Time:                  16:45:25          Log-Likelihood:             12.958
No. Observations:      132              AIC:                       -15.92
Df Residuals:          127              BIC:                       -1.503
Df Model:              4
Covariance Type:       nonrobust

```

```

=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                -11.6414      3.309      -3.518      0.001
-18.189    -5.094
WorkforceCount       -18.9933     14.044      -1.352      0.179
-46.785     8.798
PercentNegativeUsers    0.3077     0.251       1.228      0.222
-0.188     0.803
FavorOfDemocrats      17.4337     5.001       3.486      0.001
7.538     27.329
Scaled_SearchCountForDepression -0.3310     0.092      -3.608      0.000
-0.513    -0.149
=====
Omnibus:              5.838    Durbin-Watson:              0.144
Prob(Omnibus):        0.054    Jarque-Bera (JB):              6.010
Skew:                 -0.510    Prob(JB):                     0.0495
Kurtosis:              2.774    Cond. No.                     1.05e+03
=====


```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.05e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model without ground truth:



Covid Data Comparison:

Regression from Mar-2020 to Dec-2021 along with Covid Positive Cases Per Month:

Our Model:

Linear Model summary:

OLS Regression Results				
=====				
Dep. Variable:	Scaled_CrimeCount	R-squared:	0.491	
Model:	OLS	Adj. R-squared:	0.331	
Method:	Least Squares	F-statistic:	3.082	
Date:	Sun, 11 Dec 2022	Prob (F-statistic):	0.0389	
Time:	16:45:25	Log-Likelihood:	41.161	
No. Observations:	22	AIC:	-70.32	
Df Residuals:	16	BIC:	-63.78	
Df Model:	5			
Covariance Type:	nonrobust			
=====				
=====				
		coef	std err	t P> t
[0.025	0.975]			
-----				
Gender_ratio		0.3230	0.120	2.689 0.016
0.068	0.578			
Scaled_Unemployment		-0.1164	0.059	-1.961 0.067
-0.242	0.009			
Scaled MedianIncomeRate		-0.4951	0.139	-3.574 0.003

-0.789	-0.201				
WorkforceCount		0.0155	0.007	2.280	0.037
0.001	0.030				
PercentNegativeUsers		0.1907	0.326	0.585	0.567
-0.501	0.882				
FavorOfDemocrats		0.4663	0.173	2.689	0.016
0.099	0.834				
Scaled_SearchCountForDepression		0.1280	0.086	1.488	0.156
-0.054	0.310				
CovidPercPositive		-0.4447	0.238	-1.870	0.080
-0.949	0.059				
=====					
Omnibus:		3.575	Durbin-Watson:		2.116
Prob(Omnibus):		0.167	Jarque-Bera (JB):		1.457
Skew:		0.182	Prob(JB):		0.483
Kurtosis:		1.793	Cond. No.		3.38e+18
=====					

#### Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 4.04e-36. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

#### Plot for Linear Model Coefficients for Our Model:

```
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1807615713.py:1
```

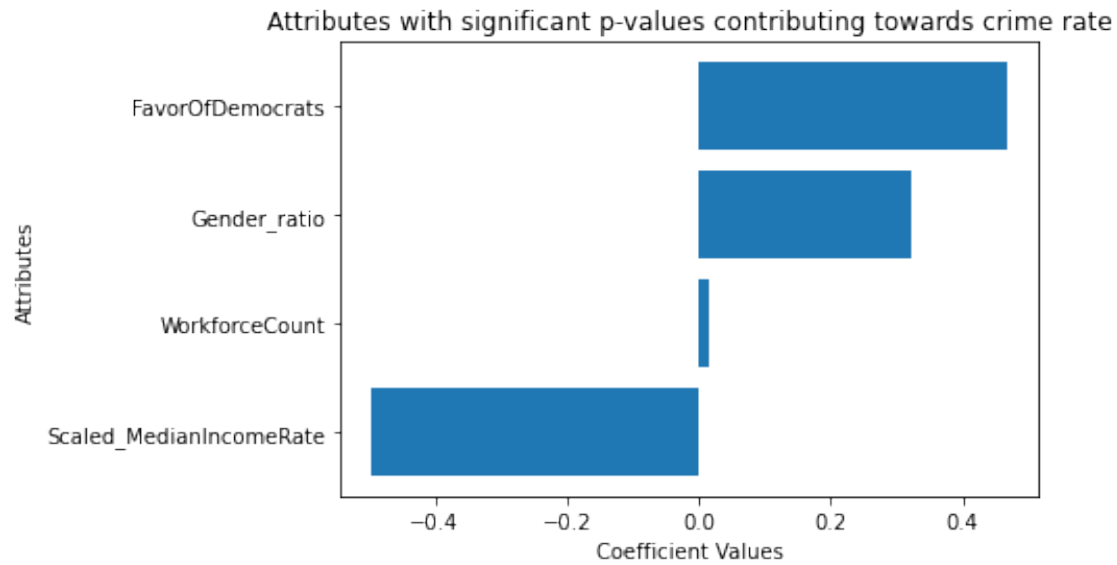
```
9: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
city_combined_dataset_covid['CovidPercPositive'] = perc
```



Regression for the same time frame without considering Covid Cases:

Our Model:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:    Scaled_CrimeCount    R-squared:        0.379
Model:            OLS                  Adj. R-squared:    0.233
Method:           Least Squares        F-statistic:       2.596
Date:             Sun, 11 Dec 2022      Prob (F-statistic): 0.0734
Time:             16:45:25              Log-Likelihood:    38.986
No. Observations: 22                   AIC:               -67.97
Df Residuals:     17                   BIC:               -62.52
Df Model:         4
Covariance Type:  nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
Gender_ratio                0.2576     0.123      2.093     0.052
-0.002      0.517
Scaled_Unemployment        -0.1220     0.063     -1.922     0.072
-0.256      0.012
Scaled_MedianIncomeRate    -0.4436     0.145     -3.051     0.007
-0.750     -0.137

```

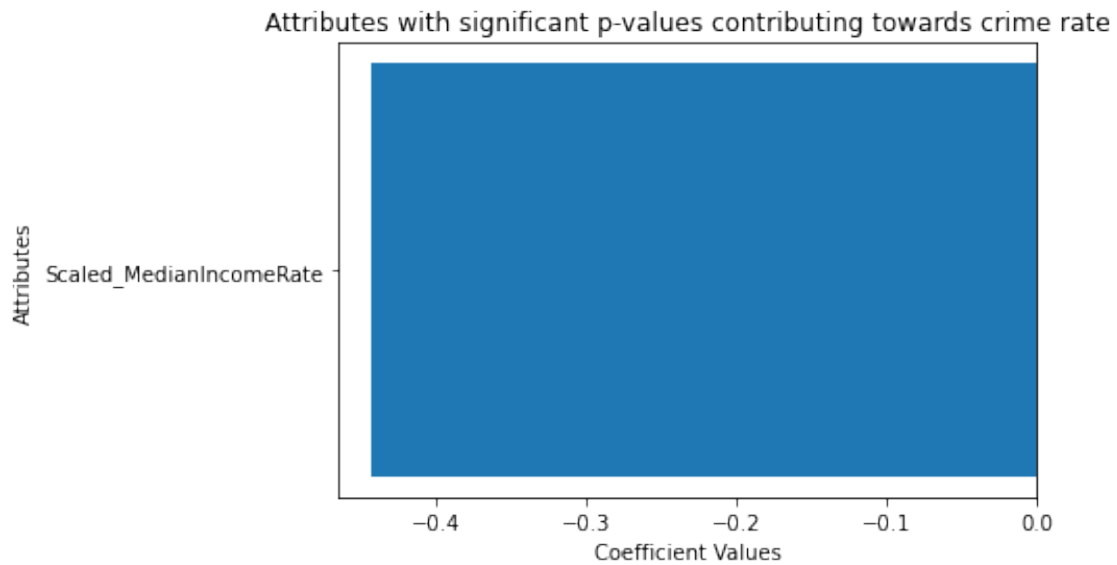


WorkforceCount	0.0119	0.007	1.701	0.107
-0.003	0.027			
PercentNegativeUsers	0.3433	0.338	1.015	0.324
-0.370	1.057			
FavorOfDemocrats	0.3719	0.178	2.093	0.052
-0.003	0.747			
Scaled_SearchCountForDepression	0.0235	0.070	0.336	0.741
-0.124	0.171			
=====				
Omnibus:	1.146	Durbin-Watson:	1.791	
Prob(Omnibus):	0.564	Jarque-Bera (JB):	0.834	
Skew:	-0.092	Prob(JB):	0.659	
Kurtosis:	2.064	Cond. No.	5.82e+18	
=====				

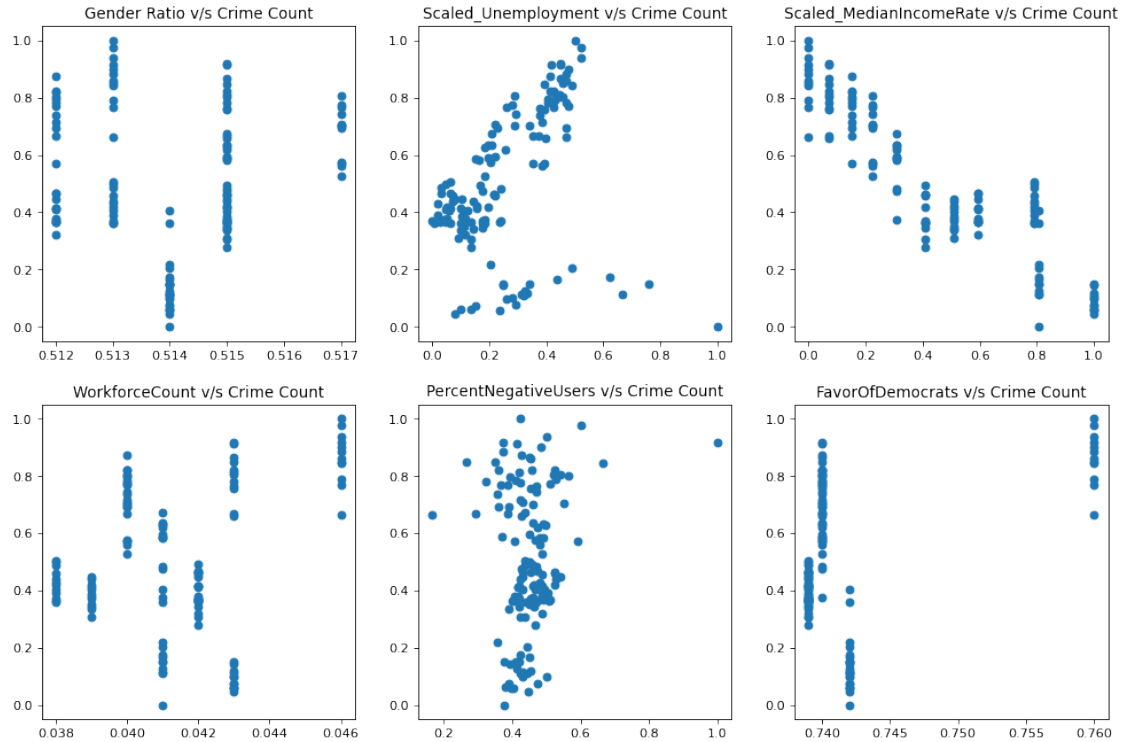
Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 1.36e-36. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Plot for Linear Model Coefficients for Our Model:



Correlation Graphs:




```

/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
1: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "g-" (-> color='g'). The keyword argument will take precedence.
    ax1.plot(city_data["MonthYear"], city_data['Scaled_Unemployment'], 'g-',
color="#6495ED")
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
2: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "b-" (-> color='b'). The keyword argument will take precedence.
    ax2.plot(city_data["MonthYear"], city_data['Scaled_CrimeCount'], 'b-',
color="#CC4F1B")

```

<Figure size 2400x1200 with 0 Axes>



NASHVILLE:

Before Covid (Mar-2020)

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.058
Model:              OLS                    Adj. R-squared:  0.009
Method:             Least Squares          F-statistic:    1.190
Date:               Sun, 11 Dec 2022        Prob (F-statistic): 0.322
Time:               16:45:26                Log-Likelihood: 16.556
No. Observations:   62                     AIC:           -25.11
Df Residuals:       58                     BIC:           -16.60
Df Model:           3
Covariance Type:    nonrobust
=====
```

```
=====
coef      std err      t      P>|t|      [0.025
```

0.975]


-----					
const	25.2928	21.064	1.201	0.235	-16.871
67.456					
Gender_ratio	-47.2977	40.717	-1.162	0.250	-128.802
34.207					
Scaled_Unemployment	-0.2374	1.015	-0.234	0.816	-2.269
1.795					
Scaled_MedianIncomeRate	-0.2129	0.152	-1.404	0.166	-0.516
0.091					
=====					
Omnibus:	13.969	Durbin-Watson:		1.196	
Prob(Omnibus):	0.001	Jarque-Bera (JB):		16.279	
Skew:	-0.966	Prob(JB):		0.000292	
Kurtosis:	4.604	Cond. No.		2.31e+03	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.31e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

# OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.346
Model:              OLS                    Adj. R-squared:  0.261
Method:             Least Squares          F-statistic:     4.081
Date:               Sun, 11 Dec 2022        Prob (F-statistic): 0.00117
Time:               16:45:26                Log-Likelihood:  27.867
No. Observations:   62                     AIC:             -39.73
Df Residuals:       54                     BIC:             -22.72
Df Model:           7
Covariance Type:    nonrobust
=====
```

```
=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
const                29.6462      25.376         1.168      0.248
-21.229      80.521
Gender_ratio         -51.5976      49.685        -1.038      0.304
-151.210      48.015
Scaled_Unemployment  -1.1478         0.913        -1.257      0.214
-2.978         0.683
Scaled_MedianIncomeRate -0.1912      0.135        -1.414      0.163
-0.462         0.080
WorkforceCount       6.1483         9.365         0.657      0.514
-12.627      24.923
PercentNegativeUsers  0.0516         0.611         0.084      0.933
-1.174         1.277
FavorOfDemocrats     -3.4526         2.498        -1.382      0.173
-8.460         1.555
Scaled_SearchCountForDepression -0.5094      0.107        -4.753      0.000
-0.724      -0.294
=====
```

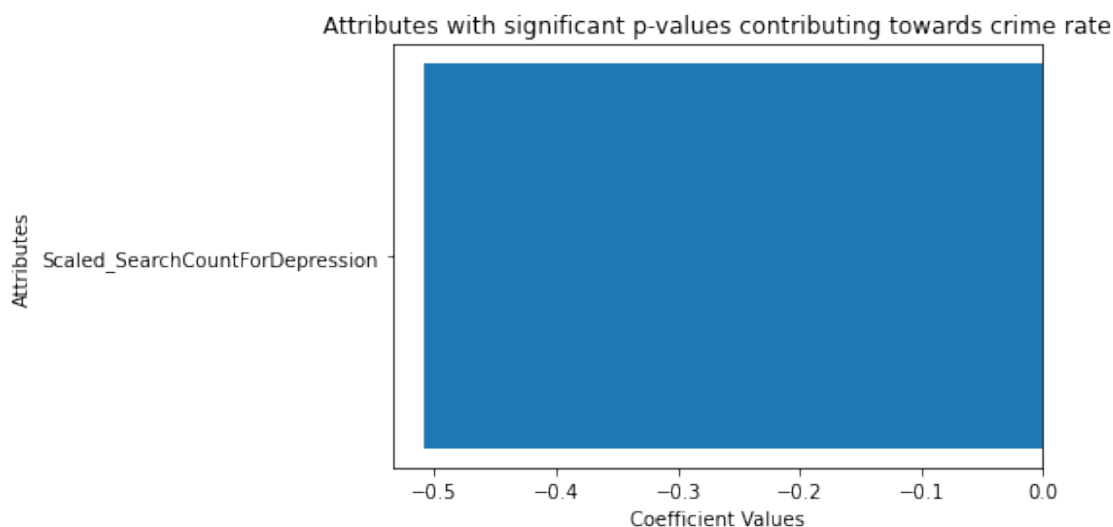
```
Omnibus:      8.402      Durbin-Watson:      1.828
Prob(Omnibus): 0.015      Jarque-Bera (JB):      7.749
Skew:         -0.760      Prob(JB):      0.0208
Kurtosis:     3.828      Cond. No.      3.97e+03
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large,  $3.97e+03$ . This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:    Scaled_CrimeCount    R-squared:        0.285
Model:            OLS                  Adj. R-squared:    0.249
Method:           Least Squares        F-statistic:       7.863
Date:            Sun, 11 Dec 2022      Prob (F-statistic): 2.18e-05
Time:            16:45:26              Log-Likelihood:    25.564
No. Observations: 84                  AIC:               -41.13
Df Residuals:    79                   BIC:               -28.97
Df Model:        4
Covariance Type: nonrobust
=====
```


```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                3.0887    0.679      4.550    0.000
1.737    4.440
WorkforceCount       -2.9960    6.354     -0.472    0.639
-15.643    9.651
=====
```

PercentNegativeUsers	0.8915	0.560	1.591	0.116
-0.224	2.007			
FavorOfDemocrats	-3.9353	0.920	-4.279	0.000
-5.766	-2.105			
Scaled_SearchCountForDepression	-0.4172	0.100	-4.188	0.000
-0.615	-0.219			
=====				
Omnibus:	5.727	Durbin-Watson:	1.328	
Prob(Omnibus):	0.057	Jarque-Bera (JB):	4.994	
Skew:	-0.551	Prob(JB):	0.0823	
Kurtosis:	3.462	Cond. No.	412.	
=====				

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Till Dec-2021

Ground Truth:

Linear Model summary:

#### OLS Regression Results

Dep. Variable:	Scaled_CrimeCount	R-squared:	0.129
----------------	-------------------	------------	-------

```

Model:                OLS      Adj. R-squared:      0.097
Method:               Least Squares    F-statistic:      3.963
Date:                 Sun, 11 Dec 2022    Prob (F-statistic): 0.0109
Time:                 16:45:26    Log-Likelihood:    17.307
No. Observations:      84    AIC:                -26.61
Df Residuals:          80    BIC:                -16.89
Df Model:              3
Covariance Type:      nonrobust

```

```

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
const              30.5696      18.361        1.665      0.100      -5.971
67.110
Gender_ratio       -57.4840      35.408       -1.623      0.108     -127.948
12.980
Scaled_Unemployment -0.1250       0.128       -0.973      0.333       -0.381
0.131
Scaled_MedianIncomeRate -0.2351      0.078       -3.025      0.003       -0.390
-0.080
=====
Omnibus:           10.873    Durbin-Watson:           1.043
Prob(Omnibus):      0.004    Jarque-Bera (JB):         11.215
Skew:              -0.756    Prob(JB):                 0.00367
Kurtosis:           3.958    Cond. No.                 2.32e+03
=====

```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.32e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Ground Truth:





Our Model:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:          Scaled_CrimeCount    R-squared:                0.319
Model:                  OLS                 Adj. R-squared:           0.257
Method:                 Least Squares       F-statistic:              5.097
Date:                   Sun, 11 Dec 2022    Prob (F-statistic):      8.82e-05
Time:                   16:45:26           Log-Likelihood:           27.655
No. Observations:       84                AIC:                     -39.31
Df Residuals:           76                BIC:                     -19.86
Df Model:                7
Covariance Type:        nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                47.7534    24.250        1.969    0.053
-0.544    96.051
Gender_ratio        -88.9294    47.761       -1.862    0.066
-184.053    6.194
Scaled_Unemployment  -0.0661    0.143       -0.461    0.646
-0.352    0.219
Scaled_MedianIncomeRate -0.1551    0.112       -1.390    0.169
-0.378    0.067
WorkforceCount       8.6963    9.930        0.876    0.384
=====
```


-11.081	28.473				
PercentNegativeUsers		0.8941	0.568	1.574	0.120
-0.237	2.025				
FavorOfDemocrats		-2.1912	1.689	-1.297	0.198
-5.555	1.173				
Scaled_SearchCountForDepression		-0.4226	0.102	-4.150	0.000
-0.625	-0.220				
=====					
Omnibus:	3.943	Durbin-Watson:		1.365	
Prob(Omnibus):	0.139	Jarque-Bera (JB):		3.213	
Skew:	-0.446	Prob(JB):		0.201	
Kurtosis:	3.351	Cond. No.		4.13e+03	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.13e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

OLS Regression Results

Dep. Variable:	Scaled_CrimeCount	R-squared:	0.285
Model:	OLS	Adj. R-squared:	0.249

```

Method:                Least Squares    F-statistic:                7.863
Date:                  Sun, 11 Dec 2022  Prob (F-statistic):       2.18e-05
Time:                  16:45:26          Log-Likelihood:            25.564
No. Observations:      84               AIC:                      -41.13
Df Residuals:          79               BIC:                      -28.97
Df Model:              4
Covariance Type:       nonrobust

```

```


=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                3.0887    0.679      4.550    0.000
1.737    4.440
WorkforceCount       -2.9960    6.354     -0.472    0.639
-15.643    9.651
PercentNegativeUsers  0.8915    0.560      1.591    0.116
-0.224    2.007
FavorOfDemocrats     -3.9353    0.920     -4.279    0.000
-5.766   -2.105
Scaled_SearchCountForDepression -0.4172    0.100     -4.188    0.000
-0.615   -0.219
=====
Omnibus:              5.727    Durbin-Watson:           1.328
Prob(Omnibus):        0.057    Jarque-Bera (JB):        4.994
Skew:                 -0.551    Prob(JB):                0.0823
Kurtosis:             3.462    Cond. No.                412.
=====

```

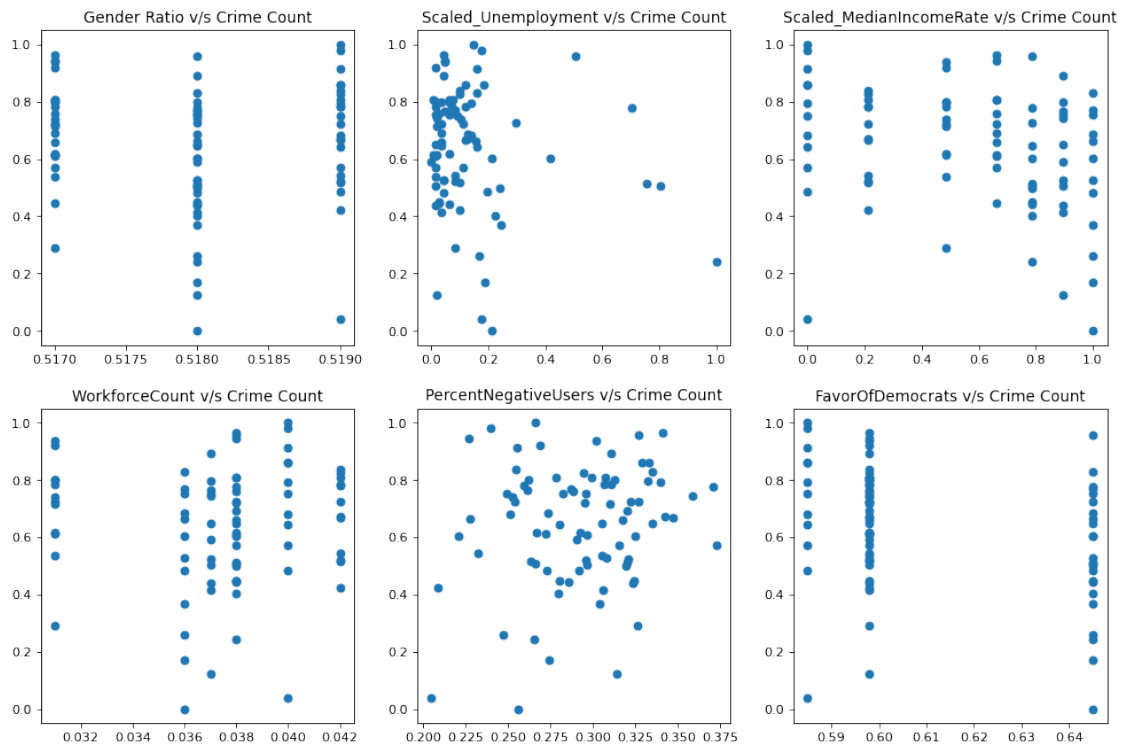
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



### Correlation Graphs:




/var/folders/j\_/kdgw7x6d25j6yr\_1c3mwp9m40000gn/T/ipykernel\_39508/1320670805.py:1  
 1: UserWarning: color is redundantly defined by the 'color' keyword argument and

```

the fmt string "g-" (-> color='g'). The keyword argument will take precedence.
ax1.plot(city_data["MonthYear"], city_data['Scaled_Unemployment'], 'g-',
color="#6495ED")
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
2: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "b-" (-> color='b'). The keyword argument will take precedence.
ax2.plot(city_data["MonthYear"], city_data['Scaled_CrimeCount'], 'b-',
color="#CC4F1B")

<Figure size 2400x1200 with 0 Axes>

```



BOSTON:

Before Covid (Mar-2020)

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:    Scaled_CrimeCount    R-squared:    0.155
Model:            OLS                  Adj. R-squared: 0.126

```

```

Method:                Least Squares    F-statistic:                5.365
Date:                  Sun, 11 Dec 2022  Prob (F-statistic):       0.00194
Time:                  16:45:27          Log-Likelihood:             36.943
No. Observations:      92              AIC:                       -65.89
Df Residuals:          88              BIC:                       -55.80
Df Model:              3
Covariance Type:       nonrobust

```

```


=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const              6.7389      5.868      1.148      0.254      -4.922
18.400
Gender_ratio      -11.0116     11.369     -0.969      0.335     -33.605
11.582
Scaled_Unemployment  -1.1918      0.396     -3.007      0.003     -1.979
-0.404
Scaled_MedianIncomeRate -0.4468      0.147     -3.037      0.003     -0.739
-0.154
=====
Omnibus:           27.005    Durbin-Watson:           1.020
Prob(Omnibus):      0.000    Jarque-Bera (JB):        55.350
Skew:              -1.095    Prob(JB):                9.57e-13
Kurtosis:           6.105    Cond. No.                 877.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount      R-squared:                0.456
Model:              OLS                    Adj. R-squared:           0.411
Method:             Least Squares          F-statistic:             10.07
Date:               Sun, 11 Dec 2022        Prob (F-statistic):       4.42e-09
Time:               16:45:27                Log-Likelihood:           57.244
No. Observations:   92                     AIC:                     -98.49
Df Residuals:       84                     BIC:                     -78.31
Df Model:           7
Covariance Type:    nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                3.6810      5.353      0.688      0.494
-6.965      14.327
Gender_ratio        -12.7624      9.638     -1.324      0.189
-31.929      6.404
Scaled_Unemployment  -0.9871      0.370     -2.670      0.009
-1.722     -0.252
Scaled_MedianIncomeRate -0.6493      0.136     -4.763      0.000
-0.920     -0.378
WorkforceCount      -7.3452      5.313     -1.383      0.170
=====
```


-17.910	3.219				
PercentNegativeUsers		-0.4246	0.251	-1.694	0.094
-0.923	0.074				
FavorOfDemocrats		5.7494	1.455	3.951	0.000
2.856	8.643				
Scaled_SearchCountForDepression		-0.3163	0.071	-4.437	0.000
-0.458	-0.175				
=====					
Omnibus:	62.383	Durbin-Watson:		1.804	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		369.691	
Skew:	-2.081	Prob(JB):		5.28e-81	
Kurtosis:	11.895	Cond. No.		1.22e+03	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.22e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

OLS Regression Results

Dep. Variable:	Scaled_CrimeCount	R-squared:	0.147
Model:	OLS	Adj. R-squared:	0.116



```

Method:                Least Squares    F-statistic:                4.704
Date:                  Sun, 11 Dec 2022  Prob (F-statistic):        0.00153
Time:                  16:45:27          Log-Likelihood:             20.606
No. Observations:      114              AIC:                       -31.21
Df Residuals:          109              BIC:                       -17.53
Df Model:               4
Covariance Type:        nonrobust

```

```


=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                2.9348      0.776      3.780      0.000
1.396      4.473
WorkforceCount       -20.1339      6.048     -3.329      0.001
-32.121     -8.147
PercentNegativeUsers  -0.5454      0.363     -1.501      0.136
-1.265      0.175
FavorOfDemocrats     -1.5521      0.945     -1.642      0.103
-3.425      0.321
Scaled_SearchCountForDepression -0.1523      0.087     -1.753      0.082
-0.324      0.020
=====
Omnibus:              13.750    Durbin-Watson:              0.626
Prob(Omnibus):         0.001    Jarque-Bera (JB):           14.911
Skew:                  -0.790    Prob(JB):                   0.000578
Kurtosis:               3.803    Cond. No.                    445.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Till Dec-2021

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:          Scaled_CrimeCount    R-squared:                0.406
Model:                  OLS                  Adj. R-squared:           0.389
Method:                 Least Squares        F-statistic:              25.02
Date:                  Sun, 11 Dec 2022      Prob (F-statistic):       2.05e-12
Time:                  16:45:27              Log-Likelihood:           41.175
No. Observations:      114                  AIC:                     -74.35
Df Residuals:          110                  BIC:                     -63.41
Df Model:               3
Covariance Type:       nonrobust
=====
```

```
=====
               coef      std err          t      P>|t|      [0.025
0.975]
-----
const               3.0473      5.229      0.583      0.561     -7.315
13.410
Gender_ratio       -4.2061     10.068     -0.418      0.677    -24.159
15.747
=====
```


Scaled_Unemployment	-0.4855	0.095	-5.132	0.000	-0.673
-0.298					
Scaled_MedianIncomeRate	-0.3119	0.053	-5.913	0.000	-0.416
-0.207					

Omnibus:	16.919	Durbin-Watson:	0.804
Prob(Omnibus):	0.000	Jarque-Bera (JB):	22.552
Skew:	-0.773	Prob(JB):	1.27e-05
Kurtosis:	4.534	Cond. No.	865.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

#### OLS Regression Results

Dep. Variable:	Scaled_CrimeCount	R-squared:	0.600
Model:	OLS	Adj. R-squared:	0.573
Method:	Least Squares	F-statistic:	22.70
Date:	Sun, 11 Dec 2022	Prob (F-statistic):	1.58e-18
Time:	16:45:27	Log-Likelihood:	63.738
No. Observations:	114	AIC:	-111.5

Df Residuals: 106 BIC: -89.59  
Df Model: 7  
Covariance Type: nonrobust


		coef	std err	t	P> t
[0.025 0.975]					
-----					
const		1.4250	5.359	0.266	0.791
-9.200	12.050				
Gender_ratio		-8.1040	9.358	-0.866	0.388
-26.656	10.448				
Scaled_Unemployment		-0.3978	0.092	-4.335	0.000
-0.580	-0.216				
Scaled_MedianIncomeRate		-0.6313	0.100	-6.305	0.000
-0.830	-0.433				
WorkforceCount		-19.3595	4.582	-4.225	0.000
-28.443	-10.276				
PercentNegativeUsers		-0.3063	0.256	-1.197	0.234
-0.813	0.201				
FavorOfDemocrats		5.9098	1.381	4.278	0.000
3.171	8.649				
Scaled_SearchCountForDepression		-0.2234	0.063	-3.547	0.001
-0.348	-0.099				
=====					
Omnibus:	43.500	Durbin-Watson:		1.314	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		133.275	
Skew:	-1.356	Prob(JB):		1.15e-29	
Kurtosis:	7.550	Cond. No.		1.29e+03	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.29e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:    Scaled_CrimeCount    R-squared:                0.147
Model:            OLS                  Adj. R-squared:           0.116
Method:           Least Squares        F-statistic:              4.704
Date:             Sun, 11 Dec 2022     Prob (F-statistic):       0.00153
Time:             16:45:27             Log-Likelihood:           20.606
No. Observations: 114                 AIC:                     -31.21
Df Residuals:     109                 BIC:                     -17.53
Df Model:         4
Covariance Type:  nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                2.9348    0.776        3.780    0.000
1.396    4.473
WorkforceCount       -20.1339    6.048       -3.329    0.001
-32.121   -8.147
PercentNegativeUsers  -0.5454    0.363       -1.501    0.136
-1.265    0.175
FavorOfDemocrats     -1.5521    0.945       -1.642    0.103
-3.425    0.321
Scaled_SearchCountForDepression -0.1523    0.087       -1.753    0.082
-0.324    0.020
=====


```

```
=====
Omnibus:                13.750    Durbin-Watson:                0.626
Prob(Omnibus):           0.001    Jarque-Bera (JB):          14.911
Skew:                   -0.790    Prob(JB):                  0.000578
Kurtosis:                3.803    Cond. No.                  445.
=====
```

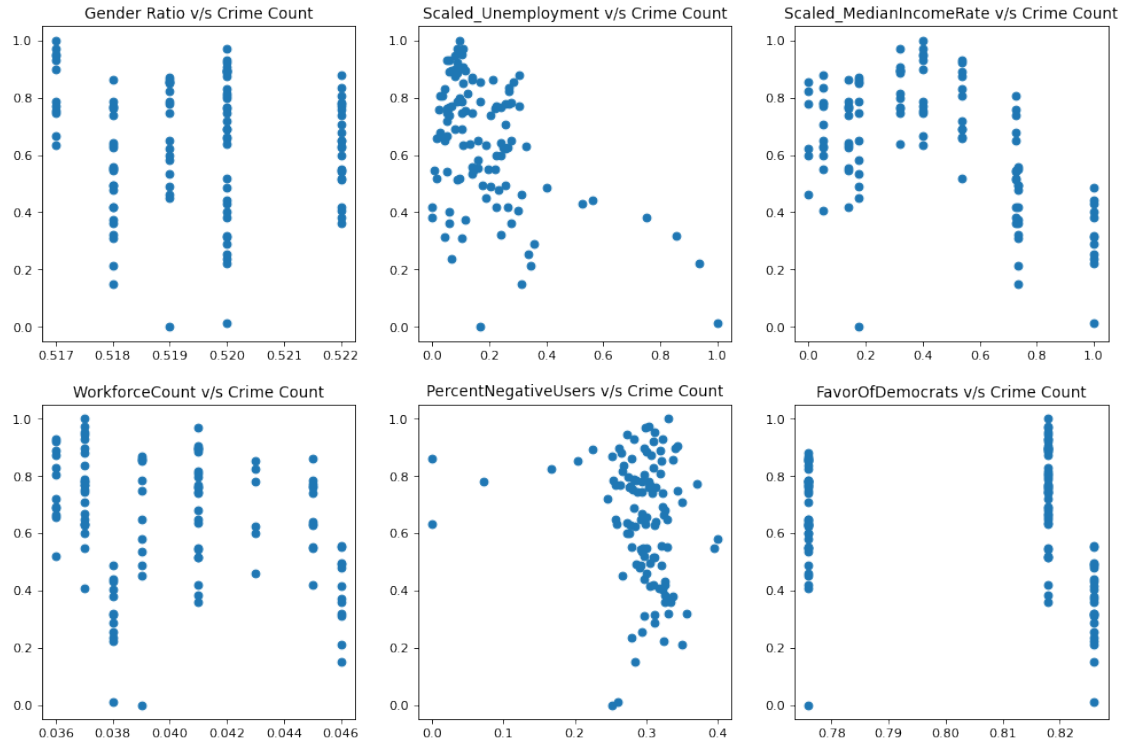
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Correlation Graphs:




```

/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
1: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "g-" (-> color='g'). The keyword argument will take precedence.
    ax1.plot(city_data["MonthYear"], city_data['Scaled_Unemployment'], 'g-',
color="#6495ED")
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
2: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "b-" (-> color='b'). The keyword argument will take precedence.
    ax2.plot(city_data["MonthYear"], city_data['Scaled_CrimeCount'], 'b-',
color="#CC4F1B")

```

<Figure size 2400x1200 with 0 Axes>



SEATTLE:

Before Covid (Mar-2020)

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.179
Model:              OLS                    Adj. R-squared:  0.156
Method:             Least Squares          F-statistic:    7.701
Date:               Sun, 11 Dec 2022        Prob (F-statistic): 0.000105
Time:               16:45:28                Log-Likelihood:  186.98
No. Observations:   110                    AIC:            -366.0
Df Residuals:       106                    BIC:            -355.2
Df Model:           3
Covariance Type:    nonrobust
=====
```

```
=====
coef      std err      t      P>|t|      [0.025
```



0.975]


const	13.7158	3.497	3.923	0.000	6.783
20.648					
Gender_ratio	-27.0115	6.995	-3.862	0.000	-40.879
-13.144					
Scaled_Unemployment	-0.0643	0.065	-0.993	0.323	-0.193
0.064					
Scaled_MedianIncomeRate	-0.1855	0.041	-4.526	0.000	-0.267
-0.104					
=====					
Omnibus:	15.992	Durbin-Watson:		1.581	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		20.839	
Skew:	-0.761	Prob(JB):		2.98e-05	
Kurtosis:	4.494	Cond. No.		2.14e+03	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.14e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

OLS Regression Results

```

=====
Dep. Variable:          Scaled_CrimeCount    R-squared:                0.269
Model:                  OLS                  Adj. R-squared:           0.219
Method:                 Least Squares        F-statistic:              5.368
Date:                   Sun, 11 Dec 2022     Prob (F-statistic):       2.95e-05
Time:                   16:45:28             Log-Likelihood:           193.39
No. Observations:       110                 AIC:                     -370.8
Df Residuals:           102                 BIC:                     -349.2
Df Model:               7
Covariance Type:        nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                        14.7662      3.506      4.212      0.000
7.812    21.720
Gender_ratio                 -30.5726      6.887     -4.439      0.000
-44.232   -16.913
Scaled_Unemployment          -0.0263      0.070     -0.375      0.709
-0.165      0.113
Scaled_MedianIncomeRate       0.0180      0.086      0.209      0.835
-0.153      0.189
WorkforceCount               26.9918     10.923      2.471      0.015
5.326    48.657
PercentNegativeUsers          0.0696      0.131      0.532      0.596
-0.190      0.329
FavorOfDemocrats             -0.3303      0.580     -0.569      0.571
-1.482      0.821
Scaled_SearchCountForDepression -0.0544      0.022     -2.430      0.017
-0.099    -0.010
=====
Omnibus:                    27.653    Durbin-Watson:              1.926
Prob(Omnibus):              0.000    Jarque-Bera (JB):           48.804
Skew:                       -1.077    Prob(JB):                   2.53e-11
Kurtosis:                   5.451    Cond. No.                   3.93e+03
=====


```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.93e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:          Scaled_CrimeCount      R-squared:                0.008
Model:                  OLS                   Adj. R-squared:           -0.024
Method:                 Least Squares         F-statistic:              0.2467
Date:                  Sun, 11 Dec 2022       Prob (F-statistic):       0.911
Time:                  16:45:28              Log-Likelihood:           133.24
No. Observations:      132                   AIC:                     -256.5
Df Residuals:          127                   BIC:                     -242.1
Df Model:               4
Covariance Type:       nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                -0.1487      0.592     -0.251     0.802
-1.319      1.022
WorkforceCount        1.2868      5.360      0.240     0.811
-9.320     11.894
PercentNegativeUsers  -0.1390      0.231     -0.602     0.548
-0.596      0.318
FavorOfDemocrats      0.4476      0.625      0.716     0.475
-0.789      1.684
=====

```

Scaled_SearchCountForDepression	0.0149	0.039	0.380	0.704
-0.062	0.092			

```
=====
```


Omnibus:	193.546	Durbin-Watson:	1.969
Prob(Omnibus):	0.000	Jarque-Bera (JB):	14902.790
Skew:	5.667	Prob(JB):	0.00
Kurtosis:	53.805	Cond. No.	935.

```
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Till Dec-2021

Ground Truth:

Linear Model summary:

OLS Regression Results

```
=====
```

```

Dep. Variable:      Scaled_CrimeCount      R-squared:      0.093
Model:              OLS                    Adj. R-squared:  0.072
Method:             Least Squares          F-statistic:     4.372
Date:              Sun, 11 Dec 2022        Prob (F-statistic): 0.00576
Time:              16:45:28                Log-Likelihood:  139.17
No. Observations:   132                    AIC:             -270.3
Df Residuals:       128                    BIC:             -258.8
Df Model:           3
Covariance Type:    nonrobust

```

```

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const              18.1398        6.188        2.931      0.004        5.895
30.385
Gender_ratio      -35.9424       12.359       -2.908      0.004       -60.397
-11.487
Scaled_Unemployment  0.1281        0.049        2.618      0.010         0.031
0.225
Scaled_MedianIncomeRate -0.1940       0.076       -2.566      0.011       -0.344
-0.044
=====
Omnibus:           173.240    Durbin-Watson:      2.156
Prob(Omnibus):      0.000    Jarque-Bera (JB):  10191.230
Skew:               4.738    Prob(JB):          0.00
Kurtosis:           44.990    Cond. No.          2.26e+03
=====


```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.26e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:          Scaled_CrimeCount      R-squared:                0.119
Model:                  OLS                   Adj. R-squared:           0.070
Method:                 Least Squares         F-statistic:              2.401
Date:                  Sun, 11 Dec 2022       Prob (F-statistic):       0.0244
Time:                  16:45:28               Log-Likelihood:           141.12
No. Observations:      132                   AIC:                     -266.2
Df Residuals:          124                   BIC:                     -243.2
Df Model:              7
Covariance Type:       nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                18.9070         6.342        2.981     0.003
6.354    31.460
Gender_ratio        -38.2771        12.504       -3.061     0.003
-63.026   -13.528
Scaled_Unemployment    0.1792         0.062        2.883     0.005
0.056     0.302
Scaled_MedianIncomeRate  0.0345         0.169        0.205     0.838
-0.299     0.368
WorkforceCount       27.5330        21.478        1.282     0.202

```


-14.978	70.044				
PercentNegativeUsers		-0.0811	0.225	-0.360	0.719
-0.527	0.364				
FavorOfDemocrats		-0.8022	0.750	-1.069	0.287
-2.287	0.683				
Scaled_SearchCountForDepression		-0.0156	0.038	-0.410	0.682
-0.091	0.060				
=====					
Omnibus:	170.283	Durbin-Watson:		2.266	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		9831.650	
Skew:	4.600	Prob(JB):		0.00	
Kurtosis:	44.266	Cond. No.		4.38e+03	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.38e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

OLS Regression Results

=====

```

Dep. Variable:    Scaled_CrimeCount    R-squared:            0.008
Model:                OLS    Adj. R-squared:        -0.024
Method:            Least Squares    F-statistic:        0.2467
Date:                Sun, 11 Dec 2022    Prob (F-statistic):    0.911
Time:                16:45:28    Log-Likelihood:    133.24
No. Observations:    132    AIC:                -256.5
Df Residuals:        127    BIC:                -242.1
Df Model:            4
Covariance Type:    nonrobust

```

```

=====
=====

```


	coef	std err	t	P> t
[0.025      0.975]				
-----				
const	-0.1487	0.592	-0.251	0.802
-1.319      1.022				
WorkforceCount	1.2868	5.360	0.240	0.811
-9.320      11.894				
PercentNegativeUsers	-0.1390	0.231	-0.602	0.548
-0.596      0.318				
FavorOfDemocrats	0.4476	0.625	0.716	0.475
-0.789      1.684				
Scaled_SearchCountForDepression	0.0149	0.039	0.380	0.704
-0.062      0.092				
=====				
Omnibus:	193.546	Durbin-Watson:	1.969	
Prob(Omnibus):	0.000	Jarque-Bera (JB):	14902.790	
Skew:	5.667	Prob(JB):	0.00	
Kurtosis:	53.805	Cond. No.	935.	
=====				

Notes:

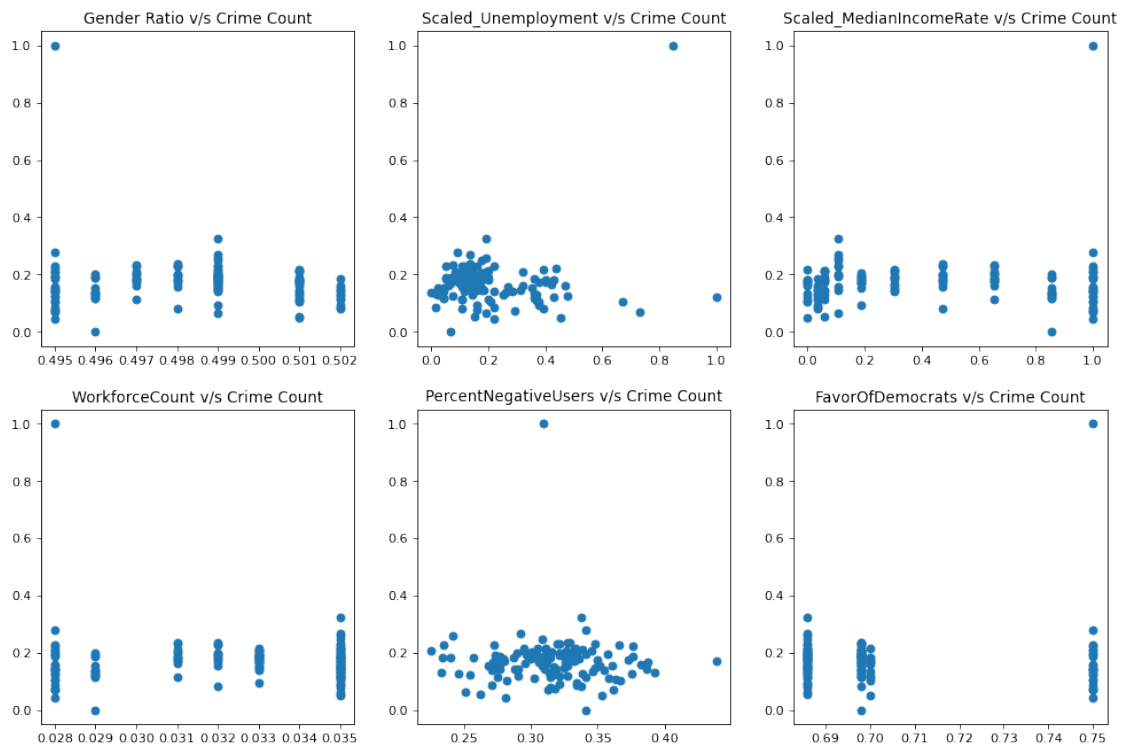
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:





Correlation Graphs:



```

/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
1: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "g-" (-> color='g'). The keyword argument will take precedence.
    ax1.plot(city_data["MonthYear"], city_data['Scaled_Unemployment'], 'g-',
color="#6495ED")
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
2: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "b-" (-> color='b'). The keyword argument will take precedence.
    ax2.plot(city_data["MonthYear"], city_data['Scaled_CrimeCount'], 'b-',
color="#CC4F1B")

```

<Figure size 2400x1200 with 0 Axes>



DENVER:

Before Covid (Mar-2020)

Ground Truth:

# Linear Model summary:

## OLS Regression Results

```

=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.804
Model:              OLS                    Adj. R-squared:  0.798
Method:             Least Squares          F-statistic:    144.8
Date:               Sun, 11 Dec 2022        Prob (F-statistic): 2.41e-37
Time:               16:45:29                Log-Likelihood:  84.786
No. Observations:   110                    AIC:            -161.6
Df Residuals:       106                    BIC:            -150.8
Df Model:           3
Covariance Type:    nonrobust
=====

```

```

=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
const                -2.9581      13.854      -0.214      0.831     -30.425
24.509
Gender_ratio          5.5403      27.692       0.200      0.842     -49.361
60.442
Scaled_Unemployment    0.3660       0.103       3.566      0.001       0.163
0.569
Scaled_MedianIncomeRate 1.1174       0.094      11.901      0.000       0.931
1.304
=====
Omnibus:              45.501      Durbin-Watson:      0.392
Prob(Omnibus):         0.000      Jarque-Bera (JB):    187.698
Skew:                  -1.327      Prob(JB):            1.75e-41
Kurtosis:              8.823      Cond. No.            3.40e+03
=====


```

### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.4e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.891
Model:              OLS                    Adj. R-squared:  0.884
Method:             Least Squares          F-statistic:    119.6
Date:               Sun, 11 Dec 2022        Prob (F-statistic): 3.04e-46
Time:               16:45:29                Log-Likelihood:  117.32
No. Observations:   110                    AIC:            -218.6
Df Residuals:       102                    BIC:            -197.0
Df Model:           7
Covariance Type:    nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                34.4739      12.006      2.871      0.005
10.660      58.288
Gender_ratio         -58.1437      23.386     -2.486      0.015
-104.530    -11.758
Scaled_Unemployment    0.4250       0.088      4.825      0.000
0.250       0.600
Scaled_MedianIncomeRate 1.2283       0.074     16.706      0.000
1.082       1.374
WorkforceCount        3.6717       3.239      1.134      0.260
```


-2.752	10.096				
PercentNegativeUsers		-0.1995	0.247	-0.806	0.422
-0.690	0.291				
FavorOfDemocrats		-7.6505	1.039	-7.361	0.000
-9.712	-5.589				
Scaled_SearchCountForDepression		-0.2229	0.046	-4.805	0.000
-0.315	-0.131				
=====					
Omnibus:	3.928	Durbin-Watson:		0.494	
Prob(Omnibus):	0.140	Jarque-Bera (JB):		2.188	
Skew:	-0.021	Prob(JB):		0.335	
Kurtosis:	2.310	Cond. No.		4.84e+03	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.84e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

OLS Regression Results

Dep. Variable:	Scaled_CrimeCount	R-squared:	0.277
Model:	OLS	Adj. R-squared:	0.254

```

Method:                Least Squares    F-statistic:                12.16
Date:                  Sun, 11 Dec 2022  Prob (F-statistic):        2.12e-08
Time:                  16:45:29          Log-Likelihood:             26.188
No. Observations:      132              AIC:                       -42.38
Df Residuals:          127              BIC:                       -27.96
Df Model:               4
Covariance Type:       nonrobust

```

```


=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                2.9093      0.682      4.268      0.000
1.560      4.258
WorkforceCount       -21.5788     6.051     -3.566      0.001
-33.553     -9.605
PercentNegativeUsers    1.5894     0.446     3.564      0.001
0.707      2.472
FavorOfDemocrats      -3.0935     0.843     -3.669      0.000
-4.762     -1.425
Scaled_SearchCountForDepression  0.0762     0.089     0.854      0.395
-0.100     0.253
=====
Omnibus:              13.018    Durbin-Watson:              0.272
Prob(Omnibus):         0.001    Jarque-Bera (JB):           13.886
Skew:                  0.771    Prob(JB):                   0.000966
Kurtosis:              3.382    Cond. No.                    472.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Till Dec-2021

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:          Scaled_CrimeCount    R-squared:                0.596
Model:                  OLS                  Adj. R-squared:           0.586
Method:                 Least Squares        F-statistic:              62.89
Date:                   Sun, 11 Dec 2022     Prob (F-statistic):       4.69e-25
Time:                   16:45:29             Log-Likelihood:           64.567
No. Observations:       132                  AIC:                      -121.1
Df Residuals:           128                  BIC:                      -109.6
Df Model:                3
Covariance Type:        nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const                -41.3537      9.626     -4.296     0.000    -60.401
-22.306
Gender_ratio          83.1072     19.225      4.323     0.000     45.067
121.147
Scaled_Unemployment   -0.3693      0.067     -5.484     0.000     -0.503
=====

```

-0.236  
 Scaled\_MedianIncomeRate 0.5486 0.071 7.724 0.000 0.408  
 0.689


```
=====
Omnibus: 3.288 Durbin-Watson: 0.248
Prob(Omnibus): 0.193 Jarque-Bera (JB): 2.970
Skew: 0.366 Prob(JB): 0.227
Kurtosis: 3.072 Cond. No. 2.02e+03
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.02e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:  Scaled_CrimeCount  R-squared: 0.807
Model:  OLS  Adj. R-squared: 0.796
Method:  Least Squares  F-statistic: 73.99
Date:  Sun, 11 Dec 2022  Prob (F-statistic): 3.08e-41
Time:  16:45:29  Log-Likelihood: 113.30
=====
```



```

No. Observations:      132    AIC:                -210.6
Df Residuals:          124    BIC:                -187.5
Df Model:              7
Covariance Type:      nonrobust

```

```

=====
=====

```


		coef	std err	t	P> t
-----					
	[0.025      0.975]				
-----					
const		-18.9525	8.401	-2.256	0.026
	-35.580      -2.325				
Gender_ratio		51.3957	16.213	3.170	0.002
	19.306      83.485				
Scaled_Unemployment		0.1716	0.073	2.357	0.020
	0.028      0.316				
Scaled_MedianIncomeRate		1.1022	0.074	14.974	0.000
	0.957      1.248				
WorkforceCount		9.4968	3.816	2.489	0.014
	1.944      17.049				
PercentNegativeUsers		-0.1267	0.257	-0.492	0.624
	-0.636      0.383				
FavorOfDemocrats		-9.6450	0.900	-10.722	0.000
	-11.425      -7.865				
Scaled_SearchCountForDepression		-0.1330	0.049	-2.715	0.008
	-0.230      -0.036				
=====					
Omnibus:		25.056	Durbin-Watson:		0.273
Prob(Omnibus):		0.000	Jarque-Bera (JB):		9.581
Skew:		0.425	Prob(JB):		0.00831
Kurtosis:		1.990	Cond. No.		3.08e+03
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.08e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:    Scaled_CrimeCount    R-squared:                0.277
Model:            OLS                  Adj. R-squared:           0.254
Method:           Least Squares        F-statistic:              12.16
Date:             Sun, 11 Dec 2022     Prob (F-statistic):       2.12e-08
Time:             16:45:29             Log-Likelihood:           26.188
No. Observations: 132                 AIC:                      -42.38
Df Residuals:     127                 BIC:                      -27.96
Df Model:         4
Covariance Type:  nonrobust
=====
```


		coef	std err	t	P> t
[0.025	0.975]				
-----					
const		2.9093	0.682	4.268	0.000
1.560	4.258				
WorkforceCount		-21.5788	6.051	-3.566	0.001
-33.553	-9.605				
PercentNegativeUsers		1.5894	0.446	3.564	0.001
0.707	2.472				
FavorOfDemocrats		-3.0935	0.843	-3.669	0.000
-4.762	-1.425				
Scaled_SearchCountForDepression		0.0762	0.089	0.854	0.395
-0.100	0.253				

Omnibus:	13.018	Durbin-Watson:	0.272
Prob(Omnibus):	0.001	Jarque-Bera (JB):	13.886
Skew:	0.771	Prob(JB):	0.000966
Kurtosis:	3.382	Cond. No.	472.

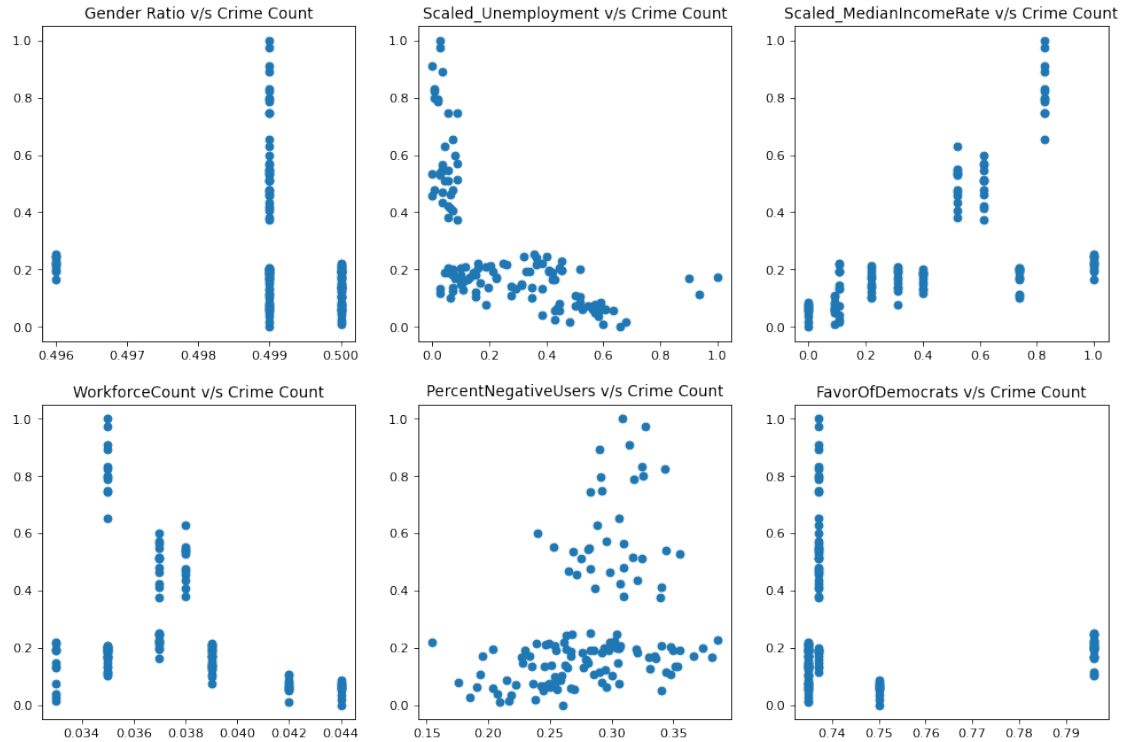
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Correlation Graphs:




```

/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
1: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "g-" (-> color='g'). The keyword argument will take precedence.
    ax1.plot(city_data["MonthYear"], city_data['Scaled_Unemployment'], 'g-',
color="#6495ED")
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
2: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "b-" (-> color='b'). The keyword argument will take precedence.
    ax2.plot(city_data["MonthYear"], city_data['Scaled_CrimeCount'], 'b-',
color="#CC4F1B")

```

<Figure size 2400x1200 with 0 Axes>



DALLAS:

Before Covid (Mar-2020)

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.527
Model:              OLS                    Adj. R-squared:  0.505
Method:             Least Squares          F-statistic:    24.10
Date:               Sun, 11 Dec 2022        Prob (F-statistic): 1.33e-10
Time:               16:45:30                Log-Likelihood:  41.919
No. Observations:   69                     AIC:            -75.84
Df Residuals:       65                     BIC:            -66.90
Df Model:           3
Covariance Type:    nonrobust
=====
```

```
=====
coef      std err      t      P>|t|      [0.025
```


0.975]

const	4.5994	3.188	1.443	0.154	-1.768
10.967					
Gender_ratio	-8.1077	6.305	-1.286	0.203	-20.700
4.485					
Scaled_Unemployment	1.1476	0.433	2.653	0.010	0.284
2.012					
Scaled_MedianIncomeRate	-0.2790	0.093	-2.999	0.004	-0.465
-0.093					
=====					
Omnibus:	8.945	Durbin-Watson:		0.873	
Prob(Omnibus):	0.011	Jarque-Bera (JB):		10.962	
Skew:	0.544	Prob(JB):		0.00416	
Kurtosis:	4.622	Cond. No.		524.	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

OLS Regression Results

```

=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.563
Model:              OLS                    Adj. R-squared:  0.513
Method:             Least Squares          F-statistic:    11.24
Date:               Sun, 11 Dec 2022       Prob (F-statistic): 4.76e-09
Time:               16:45:30              Log-Likelihood:  44.706
No. Observations:   69                   AIC:            -73.41
Df Residuals:       61                   BIC:            -55.54
Df Model:           7
Covariance Type:    nonrobust
=====

=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
const                        -5.3942      5.639      -0.957      0.343
-16.671      5.883
Gender_ratio                  6.5236      9.218       0.708      0.482
-11.909      24.956
Scaled_Unemployment           1.5814      0.512       3.086      0.003
0.557      2.606
Scaled_MedianIncomeRate      -0.5540      0.159      -3.488      0.001
-0.871      -0.236
WorkforceCount               30.1947     19.269       1.567      0.122
-8.335      68.725
PercentNegativeUsers          0.0567      0.268       0.212      0.833
-0.479      0.592
FavorOfDemocrats              3.4560      1.939       1.783      0.080
-0.421      7.333
Scaled_SearchCountForDepression 0.1059      0.123       0.859      0.394
-0.140      0.352
=====
Omnibus:                    2.781      Durbin-Watson:      0.999
Prob(Omnibus):              0.249      Jarque-Bera (JB):    2.168
Skew:                      0.197      Prob(JB):            0.338
Kurtosis:                   3.774      Cond. No.            1.75e+03
=====


```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.75e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount      R-squared:                0.375
Model:              OLS                    Adj. R-squared:           0.346
Method:             Least Squares          F-statistic:             12.91
Date:               Sun, 11 Dec 2022        Prob (F-statistic):       2.81e-08
Time:               16:45:30                Log-Likelihood:           36.864
No. Observations:   91                    AIC:                     -63.73
Df Residuals:       86                    BIC:                     -51.17
Df Model:           4
Covariance Type:    nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                3.0907      0.427      7.240      0.000
2.242      3.939
WorkforceCount       -3.9987     11.473     -0.349      0.728
-26.807     18.809
PercentNegativeUsers -0.4985      0.305     -1.633      0.106
-1.105      0.108
FavorOfDemocrats     -3.8625      0.908     -4.252      0.000
-5.668     -2.057
Scaled_SearchCountForDepression -0.1874      0.098     -1.906      0.060
=====
```




-0.383      0.008

```
=====
Omnibus:                9.059    Durbin-Watson:                0.824
Prob(Omnibus):          0.011    Jarque-Bera (JB):          9.005
Skew:                   0.638    Prob(JB):                  0.0111
Kurtosis:               3.863    Cond. No.                  831.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Till Dec-2021

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:    Scaled_CrimeCount    R-squared:                0.519
Model:            OLS                  Adj. R-squared:           0.503
Method:           Least Squares        F-statistic:              31.32
=====
```


Date: Sun, 11 Dec 2022 Prob (F-statistic): 7.99e-14  
 Time: 16:45:30 Log-Likelihood: 48.780  
 No. Observations: 91 AIC: -89.56  
 Df Residuals: 87 BIC: -79.52  
 Df Model: 3  
 Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025
0.975]					
-----					
-----					
const	4.9098	2.009	2.444	0.017	0.917
8.903					
Gender_ratio	-8.2919	3.964	-2.092	0.039	-16.170
-0.413					
Scaled_Unemployment	-0.0628	0.083	-0.758	0.451	-0.228
0.102					
Scaled_MedianIncomeRate	-0.4986	0.054	-9.288	0.000	-0.605
-0.392					
=====					
Omnibus:	0.862	Durbin-Watson:		0.976	
Prob(Omnibus):	0.650	Jarque-Bera (JB):		0.883	
Skew:	0.224	Prob(JB):		0.643	
Kurtosis:	2.822	Cond. No.		372.	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:          Scaled_CrimeCount    R-squared:                0.526
Model:                  OLS                  Adj. R-squared:           0.486
Method:                 Least Squares        F-statistic:              13.17
Date:                  Sun, 11 Dec 2022      Prob (F-statistic):       2.70e-11
Time:                  16:45:30              Log-Likelihood:           49.448
No. Observations:      91                   AIC:                     -82.90
Df Residuals:          83                   BIC:                     -62.81
Df Model:               7
Covariance Type:       nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                1.1522      4.783      0.241      0.810
-8.360      10.665
Gender_ratio        -1.9348      7.844     -0.247      0.806
-17.536     13.667
Scaled_Unemployment -0.0528      0.098     -0.537      0.593
-0.248      0.143
Scaled_MedianIncomeRate -0.5992      0.162     -3.701      0.000
-0.921     -0.277
WorkforceCount      20.0799     18.777      1.069      0.288

```


-17.268	57.427				
PercentNegativeUsers		-0.0286	0.288	-0.099	0.921
-0.601	0.544				
FavorOfDemocrats		0.3751	1.605	0.234	0.816
-2.817	3.567				
Scaled_SearchCountForDepression		0.0088	0.095	0.092	0.927
-0.181	0.198				
=====					
Omnibus:	0.604	Durbin-Watson:		0.995	
Prob(Omnibus):	0.739	Jarque-Bera (JB):		0.741	
Skew:	0.159	Prob(JB):		0.690	
Kurtosis:	2.692	Cond. No.		1.95e+03	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.95e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

OLS Regression Results

Dep. Variable:	Scaled_CrimeCount	R-squared:	0.375
----------------	-------------------	------------	-------

```

Model:                OLS      Adj. R-squared:        0.346
Method:               Least Squares    F-statistic:        12.91
Date:                 Sun, 11 Dec 2022    Prob (F-statistic):    2.81e-08
Time:                 16:45:30    Log-Likelihood:        36.864
No. Observations:      91    AIC:                -63.73
Df Residuals:          86    BIC:                -51.17
Df Model:              4
Covariance Type:      nonrobust

```

```

=====
=====


```

	coef	std err	t	P> t
[0.025      0.975]				
-----				
const	3.0907	0.427	7.240	0.000
2.242      3.939				
WorkforceCount	-3.9987	11.473	-0.349	0.728
-26.807      18.809				
PercentNegativeUsers	-0.4985	0.305	-1.633	0.106
-1.105      0.108				
FavorOfDemocrats	-3.8625	0.908	-4.252	0.000
-5.668      -2.057				
Scaled_SearchCountForDepression	-0.1874	0.098	-1.906	0.060
-0.383      0.008				
=====				
Omnibus:	9.059	Durbin-Watson:	0.824	
Prob(Omnibus):	0.011	Jarque-Bera (JB):	9.005	
Skew:	0.638	Prob(JB):	0.0111	
Kurtosis:	3.863	Cond. No.	831.	
=====				

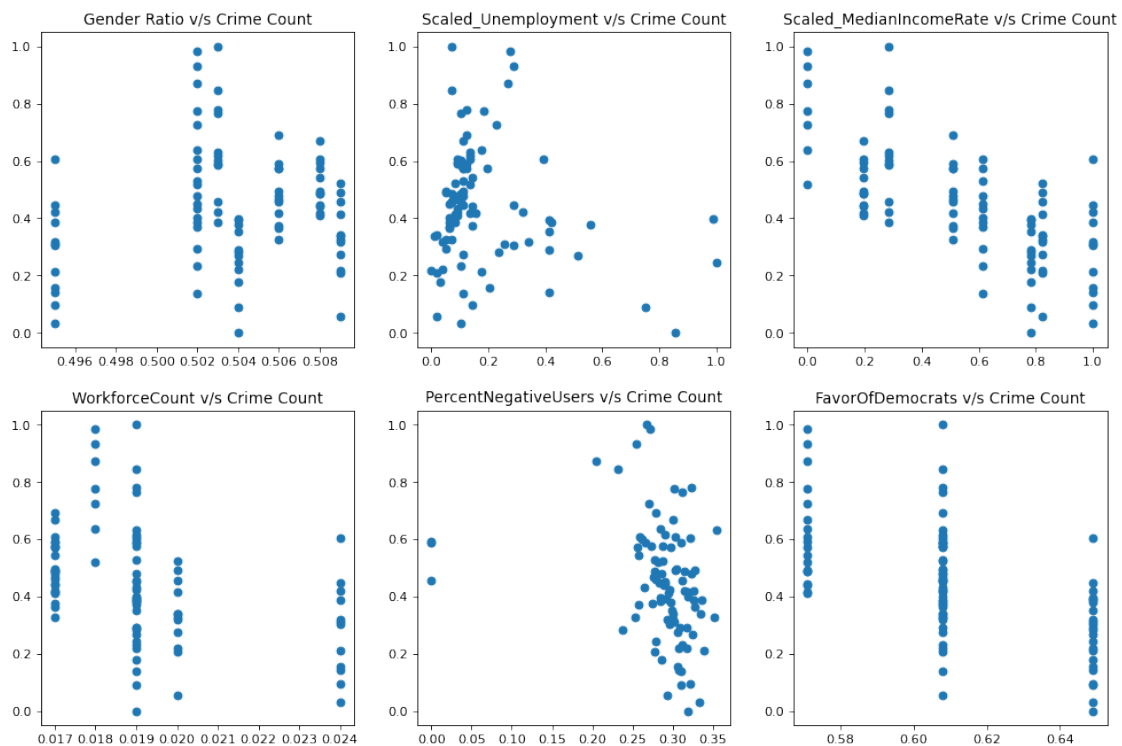
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:




### Correlation Graphs:



```

/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
1: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "g-" (-> color='g'). The keyword argument will take precedence.
    ax1.plot(city_data["MonthYear"], city_data['Scaled_Unemployment'], 'g-',
color="#6495ED")
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
2: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "b-" (-> color='b'). The keyword argument will take precedence.
    ax2.plot(city_data["MonthYear"], city_data['Scaled_CrimeCount'], 'b-',
color="#CC4F1B")
<Figure size 2400x1200 with 0 Axes>

```



NEW\_ORLEANS:

Before Covid (Mar-2020)

Ground Truth:

Linear Model summary:

OLS Regression Results

=====

```

Dep. Variable:      Scaled_CrimeCount      R-squared:      0.174
Model:              OLS                    Adj. R-squared:  0.150
Method:             Least Squares          F-statistic:     7.432
Date:              Sun, 11 Dec 2022        Prob (F-statistic): 0.000145
Time:              16:45:31                Log-Likelihood:  58.774
No. Observations:  110                    AIC:            -109.5
Df Residuals:      106                    BIC:            -98.75
Df Model:          3
Covariance Type:   nonrobust

```

```

=====
=====

```

	coef	std err	t	P> t	[0.025
0.975]					
const	17.0589	4.420	3.859	0.000	8.296
25.822					
Gender_ratio	-31.5360	8.437	-3.738	0.000	-48.262
-14.810					
Scaled_Unemployment	-0.2596	0.314	-0.826	0.411	-0.883
0.364					
Scaled_MedianIncomeRate	0.3305	0.080	4.118	0.000	0.171
0.490					

```

=====
Omnibus:          59.259    Durbin-Watson:      0.799
Prob(Omnibus):    0.000    Jarque-Bera (JB):  287.392
Skew:             -1.756    Prob(JB):          3.92e-63
Kurtosis:         10.097    Cond. No.          802.
=====


```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:





Our Model:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:    Scaled_CrimeCount    R-squared:            0.309
Model:            OLS                  Adj. R-squared:       0.262
Method:           Least Squares        F-statistic:          6.526
Date:             Sun, 11 Dec 2022      Prob (F-statistic):   2.29e-06
Time:             16:45:31              Log-Likelihood:       68.629
No. Observations: 110                  AIC:                  -121.3
Df Residuals:     102                   BIC:                  -99.65
Df Model:         7
Covariance Type:  nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                -8.4221      7.704      -1.093    0.277
-23.703      6.859
Gender_ratio         11.3918     12.625       0.902    0.369
-13.649     36.433
Scaled_Unemployment  -0.1678      0.307      -0.546    0.586
-0.777      0.442
Scaled_MedianIncomeRate -0.0557     0.125     -0.444    0.658
-0.304      0.193
WorkforceCount       53.3968     13.820       3.864    0.000
```


25.985	80.809				
PercentNegativeUsers		-0.0741	0.124	-0.597	0.552
-0.320	0.172				
FavorOfDemocrats		0.4838	4.173	0.116	0.908
-7.793	8.761				
Scaled_SearchCountForDepression		-0.0944	0.105	-0.903	0.369
-0.302	0.113				
=====					
Omnibus:	63.824	Durbin-Watson:	1.000		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	381.243		
Skew:	-1.827	Prob(JB):	1.64e-83		
Kurtosis:	11.356	Cond. No.	2.14e+03		
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.14e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

OLS Regression Results

```

=====
Dep. Variable:    Scaled_CrimeCount    R-squared:            0.346
Model:                OLS    Adj. R-squared:        0.325
Method:            Least Squares    F-statistic:         16.79
Date:                Sun, 11 Dec 2022    Prob (F-statistic):   4.54e-11
Time:                16:45:31    Log-Likelihood:       82.423
No. Observations:    132    AIC:                 -154.8
Df Residuals:        127    BIC:                 -140.4
Df Model:            4
Covariance Type:    nonrobust
=====

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                -1.1539      1.501      -0.769      0.443
-4.123      1.816
WorkforceCount        43.2886      7.716       5.611      0.000
28.021      58.556
PercentNegativeUsers   -0.0035      0.115     -0.030      0.976
-0.231      0.224
FavorOfDemocrats      -0.6496      1.492     -0.435      0.664
-3.602      2.303
Scaled_SearchCountForDepression  0.0483      0.077       0.630      0.530
-0.103      0.200
=====


Omnibus:            49.530    Durbin-Watson:           1.008
Prob(Omnibus):       0.000    Jarque-Bera (JB):        190.272
Skew:                -1.297    Prob(JB):                 4.82e-42
Kurtosis:            8.279    Cond. No.                  944.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Till Dec-2021

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount    R-squared:                0.190
Model:              OLS                  Adj. R-squared:           0.171
Method:             Least Squares        F-statistic:              10.01
Date:               Sun, 11 Dec 2022     Prob (F-statistic):       5.68e-06
Time:               16:45:31             Log-Likelihood:           68.319
No. Observations:   132                  AIC:                      -128.6
Df Residuals:       128                  BIC:                      -117.1
Df Model:           3
Covariance Type:    nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const                16.1179      3.407      4.731      0.000      9.377
22.858
Gender_ratio         -29.6911      6.558     -4.528      0.000     -42.667
-16.716
=====
```


Scaled_Unemployment	-0.2674	0.084	-3.171	0.002	-0.434
-0.101					
Scaled_MedianIncomeRate	0.2052	0.069	2.953	0.004	0.068
0.343					

Omnibus:	33.459	Durbin-Watson:	0.765
Prob(Omnibus):	0.000	Jarque-Bera (JB):	83.602
Skew:	-0.984	Prob(JB):	7.02e-19
Kurtosis:	6.366	Cond. No.	699.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

#### OLS Regression Results

Dep. Variable:	Scaled_CrimeCount	R-squared:	0.367
Model:	OLS	Adj. R-squared:	0.331
Method:	Least Squares	F-statistic:	10.26
Date:	Sun, 11 Dec 2022	Prob (F-statistic):	4.28e-10
Time:	16:45:31	Log-Likelihood:	84.559
No. Observations:	132	AIC:	-153.1

Df Residuals: 124 BIC: -130.1  
Df Model: 7  
Covariance Type: nonrobust


		coef	std err	t	P> t
[0.025 0.975]					
-----					
const		-10.1459	6.825	-1.487	0.140
-23.655	3.363				
Gender_ratio		16.1323	10.553	1.529	0.129
-4.755	37.020				
Scaled_Unemployment		-0.0304	0.089	-0.343	0.732
-0.206	0.145				
Scaled_MedianIncomeRate		-0.0911	0.114	-0.795	0.428
-0.318	0.136				
WorkforceCount		55.2672	13.302	4.155	0.000
28.940	81.595				
PercentNegativeUsers		-0.0470	0.117	-0.403	0.688
-0.278	0.184				
FavorOfDemocrats		-0.6228	3.539	-0.176	0.861
-7.627	6.381				
Scaled_SearchCountForDepression		-0.0381	0.090	-0.422	0.674
-0.217	0.141				
=====					
Omnibus:	59.073	Durbin-Watson:		1.035	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		291.185	
Skew:	-1.484	Prob(JB):		5.89e-64	
Kurtosis:	9.643	Cond. No.		2.25e+03	
=====					

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.25e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:    Scaled_CrimeCount    R-squared:                0.346
Model:            OLS                  Adj. R-squared:           0.325
Method:           Least Squares        F-statistic:              16.79
Date:             Sun, 11 Dec 2022     Prob (F-statistic):       4.54e-11
Time:             16:45:31             Log-Likelihood:           82.423
No. Observations: 132                 AIC:                     -154.8
Df Residuals:     127                 BIC:                     -140.4
Df Model:         4
Covariance Type:  nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                -1.1539      1.501     -0.769     0.443
-4.123      1.816
WorkforceCount       43.2886      7.716      5.611     0.000
28.021     58.556
PercentNegativeUsers -0.0035      0.115     -0.030     0.976
-0.231      0.224
FavorOfDemocrats    -0.6496      1.492     -0.435     0.664
-3.602      2.303
=====
```

Scaled_SearchCountForDepression	0.0483	0.077	0.630	0.530
-0.103	0.200			

---


Omnibus:	49.530	Durbin-Watson:	1.008
Prob(Omnibus):	0.000	Jarque-Bera (JB):	190.272
Skew:	-1.297	Prob(JB):	4.82e-42
Kurtosis:	8.279	Cond. No.	944.

---

Notes:

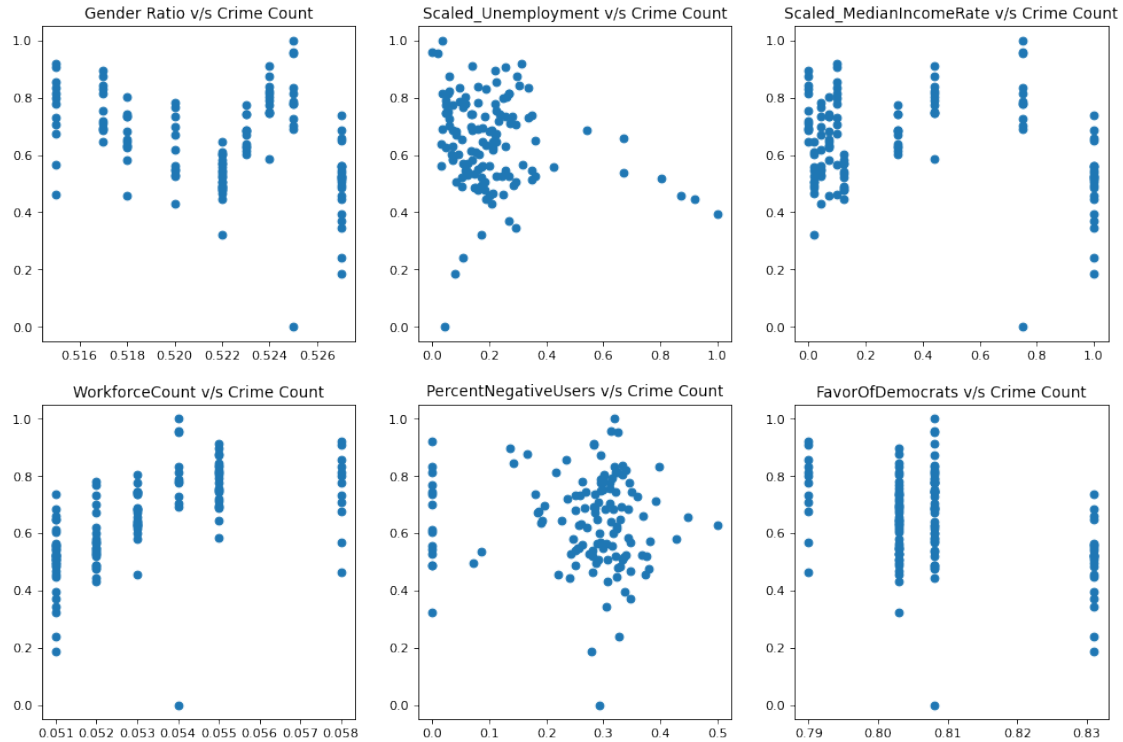
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Correlation Graphs:






```

/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
1: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "g-" (-> color='g'). The keyword argument will take precedence.
    ax1.plot(city_data["MonthYear"], city_data['Scaled_Unemployment'], 'g-',
color="#6495ED")
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
2: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "b-" (-> color='b'). The keyword argument will take precedence.
    ax2.plot(city_data["MonthYear"], city_data['Scaled_CrimeCount'], 'b-',
color="#CC4F1B")

```

<Figure size 2400x1200 with 0 Axes>



INDIANAPOLIS:

Before Covid (Mar-2020)

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.566
Model:              OLS                    Adj. R-squared:  0.554
Method:             Least Squares          F-statistic:    46.13
Date:               Sun, 11 Dec 2022        Prob (F-statistic): 3.71e-19
Time:               16:45:32                Log-Likelihood:  80.659
No. Observations:   110                    AIC:            -153.3
Df Residuals:       106                    BIC:            -142.5
Df Model:           3
Covariance Type:    nonrobust
=====
```

```
=====
coef      std err      t      P>|t|      [0.025
```


0.975]

const	-12.9353	2.881	-4.490	0.000	-18.647
-7.224					
Gender_ratio	26.1183	5.583	4.678	0.000	15.050
37.186					
Scaled_Unemployment	-0.0185	0.111	-0.166	0.868	-0.239
0.202					
Scaled_MedianIncomeRate	-0.6882	0.124	-5.550	0.000	-0.934
-0.442					
=====					
Omnibus:	3.462	Durbin-Watson:		0.711	
Prob(Omnibus):	0.177	Jarque-Bera (JB):		3.221	
Skew:	-0.211	Prob(JB):		0.200	
Kurtosis:	3.725	Cond. No.		668.	
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

OLS Regression Results

```

=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.763
Model:              OLS                    Adj. R-squared:  0.747
Method:             Least Squares          F-statistic:     46.87
Date:               Sun, 11 Dec 2022        Prob (F-statistic): 4.20e-29
Time:               16:45:32                Log-Likelihood:  113.86
No. Observations:   110                    AIC:             -211.7
Df Residuals:       102                    BIC:             -190.1
Df Model:           7
Covariance Type:    nonrobust
=====

=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
const                        -9.9486      2.314      -4.299      0.000
-14.538      -5.359
Gender_ratio                  17.3994      4.714       3.691      0.000
8.050      26.749
Scaled_Unemployment          -0.5442      0.119      -4.584      0.000
-0.780      -0.309
Scaled_MedianIncomeRate      -0.7472      0.101      -7.368      0.000
-0.948      -0.546
WorkforceCount               16.5675      4.442       3.730      0.000
7.757      25.378
PercentNegativeUsers          0.1508      0.130       1.157      0.250
-0.108      0.409
FavorOfDemocrats             2.0264      0.746       2.718      0.008
0.547      3.505
Scaled_SearchCountForDepression -0.3757      0.048      -7.860      0.000
-0.471      -0.281
=====


Omnibus:                    2.641      Durbin-Watson:      1.230
Prob(Omnibus):              0.267      Jarque-Bera (JB):    2.280
Skew:                       -0.149      Prob(JB):            0.320
Kurtosis:                   3.639      Cond. No.            975.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:          Scaled_CrimeCount      R-squared:                0.160
Model:                  OLS                   Adj. R-squared:           0.133
Method:                 Least Squares         F-statistic:              6.042
Date:                   Sun, 11 Dec 2022       Prob (F-statistic):      0.000175
Time:                   16:45:32              Log-Likelihood:          45.280
No. Observations:      132                   AIC:                     -80.56
Df Residuals:          127                   BIC:                     -66.15
Df Model:               4
Covariance Type:       nonrobust
=====

```


		coef	std err	t	P> t
[0.025	0.975]				
const		0.7992	0.456	1.752	0.082
-0.104	1.702				
WorkforceCount		12.5342	5.315	2.358	0.020
2.017	23.051				
PercentNegativeUsers		-0.4293	0.237	-1.813	0.072
-0.898	0.039				
FavorOfDemocrats		-1.1609	0.876	-1.325	0.188
-2.895	0.573				
Scaled_SearchCountForDepression		-0.2289	0.087	-2.637	0.009
-0.401	-0.057				

```
=====
Omnibus:                    5.261    Durbin-Watson:                0.360
Prob(Omnibus):              0.072    Jarque-Bera (JB):            4.935
Skew:                      0.358    Prob(JB):                    0.0848
Kurtosis:                  3.621    Cond. No.                    457.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Till Dec-2021

Ground Truth:

Linear Model summary:

#### OLS Regression Results

```
=====
Dep. Variable:    Scaled_CrimeCount    R-squared:                0.552
Model:            OLS                  Adj. R-squared:            0.542
Method:           Least Squares        F-statistic:              52.67
Date:             Sun, 11 Dec 2022     Prob (F-statistic):       3.06e-22
Time:             16:45:32             Log-Likelihood:           86.845
No. Observations: 132                 AIC:                     -165.7
Df Residuals:     128                 BIC:                     -154.2
=====
```

```


Df Model:                3
Covariance Type:         nonrobust
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const          -3.8825      2.635      -1.473      0.143     -9.097
1.332
Gender_ratio     8.4422      5.083       1.661      0.099     -1.615
18.499
Scaled_Unemployment  0.0708      0.058       1.219      0.225     -0.044
0.186
Scaled_MedianIncomeRate -0.4305      0.049     -8.788      0.000     -0.527
-0.334
=====
Omnibus:          6.143    Durbin-Watson:          0.555
Prob(Omnibus):    0.046    Jarque-Bera (JB):          7.404
Skew:             0.274    Prob(JB):              0.0247
Kurtosis:         4.023    Cond. No.              635.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Ground Truth:



Our Model:

Linear Model summary:

# OLS Regression Results

```
=====
Dep. Variable:      Scaled_CrimeCount      R-squared:      0.799
Model:              OLS                    Adj. R-squared:  0.787
Method:             Least Squares          F-statistic:    70.26
Date:               Sun, 11 Dec 2022       Prob (F-statistic): 3.94e-40
Time:               16:45:32              Log-Likelihood:  139.56
No. Observations:   132                   AIC:            -263.1
Df Residuals:       124                   BIC:            -240.1
Df Model:           7
Covariance Type:    nonrobust
=====
```

```
=====
                                coef      std err          t      P>|t|
-----
[0.025      0.975]
-----
const                                -9.5504      1.861      -5.133      0.000
-13.233      -5.868
Gender_ratio                        17.0663      3.548       4.811      0.000
10.045      24.088
Scaled_Unemployment                 -0.3287      0.059      -5.598      0.000
-0.445      -0.212
Scaled_MedianIncomeRate             -0.6004      0.040     -15.035      0.000
-0.679      -0.521
WorkforceCount                     13.4592      2.671       5.039      0.000
8.173      18.745
PercentNegativeUsers                 0.1383      0.121       1.144      0.255
-0.101      0.378
FavorOfDemocrats                    1.6635      0.538       3.091      0.002
0.598      2.729
Scaled_SearchCountForDepression     -0.3606      0.044     -8.159      0.000
-0.448      -0.273
=====
```


```
Omnibus:      6.452      Durbin-Watson:      1.211
Prob(Omnibus): 0.040      Jarque-Bera (JB):      7.442
Skew:         -0.315      Prob(JB):      0.0242
Kurtosis:     3.978      Cond. No.      789.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



Plot for Linear Model Coefficients for Our Model:



Our Model without ground truth:

Linear Model summary:

#### OLS Regression Results

```

=====
Dep. Variable:    Scaled_CrimeCount    R-squared:        0.160
Model:            OLS                  Adj. R-squared:    0.133
Method:           Least Squares        F-statistic:       6.042
Date:             Sun, 11 Dec 2022     Prob (F-statistic): 0.000175
Time:             16:45:32             Log-Likelihood:    45.280
No. Observations: 132                 AIC:               -80.56
Df Residuals:     127                 BIC:               -66.15
Df Model:         4
Covariance Type:  nonrobust
=====

```


		coef	std err	t	P> t
[0.025	0.975]				
-----					
const		0.7992	0.456	1.752	0.082
-0.104	1.702				
WorkforceCount		12.5342	5.315	2.358	0.020
2.017	23.051				
PercentNegativeUsers		-0.4293	0.237	-1.813	0.072
-0.898	0.039				

FavorOfDemocrats	-1.1609	0.876	-1.325	0.188
-2.895	0.573			
Scaled_SearchCountForDepression	-0.2289	0.087	-2.637	0.009
-0.401	-0.057			
=====				
Omnibus:	5.261	Durbin-Watson:	0.360	
Prob(Omnibus):	0.072	Jarque-Bera (JB):	4.935	
Skew:	0.358	Prob(JB):	0.0848	
Kurtosis:	3.621	Cond. No.	457.	
=====				

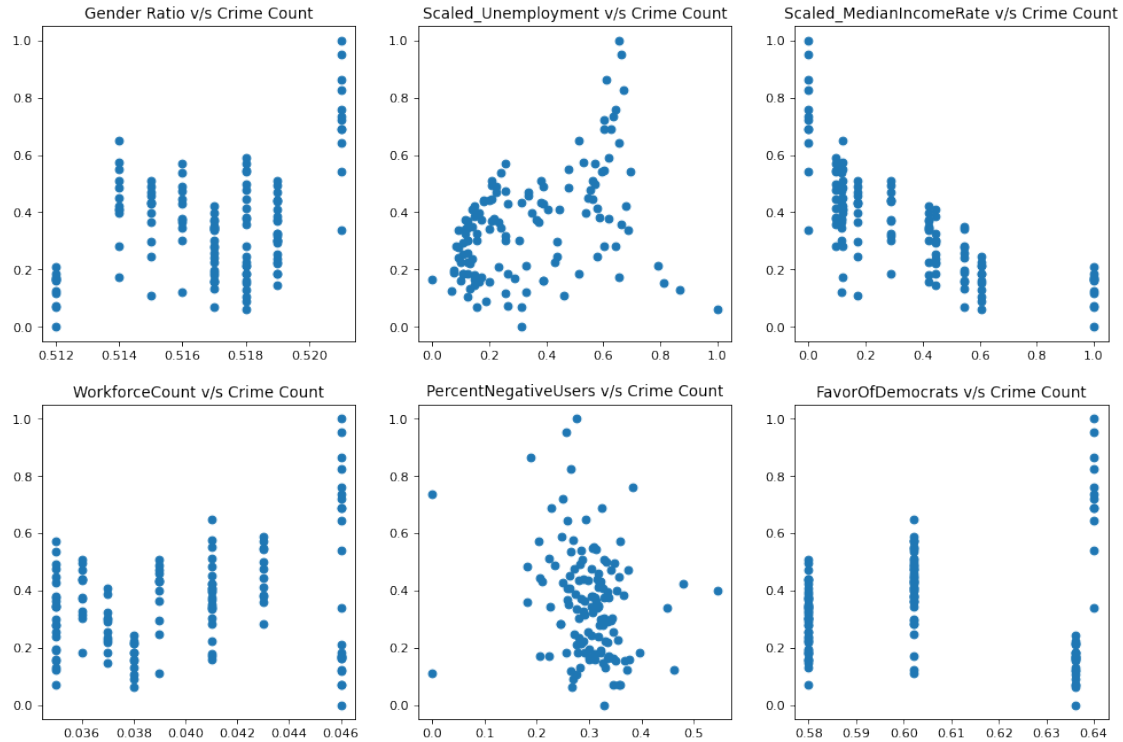
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Plot for Linear Model Coefficients for Our Model without ground truth:



Correlation Graphs:




```


/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
1: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "g-" (-> color='g'). The keyword argument will take precedence.
    ax1.plot(city_data["MonthYear"], city_data['Scaled_Unemployment'], 'g-',
color="#6495ED")
/var/folders/j_/kdgw7x6d25j6yr_1c3mwp9m40000gn/T/ipykernel_39508/1320670805.py:1
2: UserWarning: color is redundantly defined by the 'color' keyword argument and
the fmt string "b-" (-> color='b'). The keyword argument will take precedence.
    ax2.plot(city_data["MonthYear"], city_data['Scaled_CrimeCount'], 'b-',
color="#CC4F1B")

```


<Figure size 2400x1200 with 0 Axes>



Before Covid (Mar-2020)



Till Dec-2021



[25]: *# USA Cities Population Pie Chart - Cities chosen for analysis for this project*

```
USA = 331900000
```

```
total = 882039 + 654776 + 8468000 + 692587 + 376971 + 1288000 + 3849000 +  
↪ 733919 + 711463 + 2697000
```

```
def get_pop_perc(pop):  
#     return round(float(pop / USA) * 100, 2)  
    return pop/total
```

```
inidanapolis = get_pop_perc(882039)  
boston = get_pop_perc(654776)  
nyc = get_pop_perc(8468000)  
nashville = get_pop_perc(692587)  
new_orleans = get_pop_perc(376971)  
dallas = get_pop_perc(1288000)  
la = get_pop_perc(3849000)
```


```

seattle = get_pop_perc(733919)
denver = get_pop_perc(711463)
chicago = get_pop_perc(2697000)

# others = float(100 - (inidanapolis + boston + nyc + nashville + new_orleans +
↳dallas + la + seattle + denver + chicago))

figure(figsize=(8, 7), dpi=120)
labels = ["Inidanapolis", "Boston", "NYC", "Nashville", "New Orleans",
↳"Dallas", "LA", "Seattle", "Denver", "Chicago"]
plt.pie([inidanapolis, boston, nyc, nashville, new_orleans, dallas, la,
↳seattle, denver, chicago], labels=labels)
plt.legend(title = "USA Cities Population")
plt.savefig('CitySelPopWise.png', transparent=True)
plt.show()

```



```
[26]: # Playing around with FamaMacBeth Algo, with Newey-West adjustment

# import pandas as pd
# from linearmodels import FamaMacBeth # import package
# ## 'panel' is a multi-index Pandas panel dataframe (stock code - date)
# panel = pd.read_csv('~Downloads/Sample.csv')
# panel['date'] = pd.to_datetime(panel['Time'], format="%b-%y") # date should
# ↳ be `datetime` format
# panel = panel.set_index(['City', 'date']) # multi-index

# print(panel)

# mod = FamaMacBeth.from_formula('Output ~ 1 + Input1 + Input2 + Input3',
# ↳ data=panel)
# ## `bandwidth` is the lagged number of Newey-West, normally, bandwidth = 4(T/
# ↳ 100) ^ (2/9)
# ## Remove all Settings in parentheses if newey-West adjustments are not
# ↳ required.
# res = mod.fit(cov_type= 'kernel',debiased = False, bandwidth = 3)
# print(res.summary)
```

```
[27]: # Barebone implementation of FamaMacBeth

# def ols_coef(x,formula):
#     return smf.ols(formula,data=x).fit().params
# panel = pd.read_csv('~Downloads/Sample.csv')
# panel['date'] = pd.to_datetime(panel['Time'], format="%b-%y")
# gamma = (panel.groupby('date').apply(ols_coef,'Output ~ 1 + Input1 + Input2 +
# ↳ Input3'))
# print(gamma.head())

# def fm_summary(p):
#     s = p.describe().T
#     s['std_error'] = s['std']/np.sqrt(s['count'])
#     s['tstat'] = s['mean']/s['std_error']
#     return s[['mean','std_error','tstat']]
# print(fm_summary(gamma))
```

```
[28]: # Chicago Covid Comparison

# chicago_covid = pd.read_csv('~Documents/Practice/Chicago_Covid.csv',
# ↳ header=None)
# chicago_covid.columns = ['Perc', 'MonthYear']
# chicago_covid['MonthYear'] = pd.to_datetime(chicago_covid['MonthYear'],
# ↳ format="%b-%Y")
# chicago_covid = chicago_covid.sort_values(by=['MonthYear'])
```



```

# # chicago_covid    # Mar2020 to Nov2022

# chicago_dataset = pd.read_csv('~\Documents\Practice/
    ↳chicago_city_final_dataset_v2.csv')
# chicago_dataset['MonthYear'] = pd.to_datetime(chicago_dataset['MonthYear'],
    ↳format="%b-%y")

# chicago_combined_dataset_covid = chicago_dataset[chicago_dataset['MonthYear']
    ↳>= '2020-03-01']

# perc = []
# for index, row in chicago_combined_dataset_covid.iterrows():
#     month_year = row['MonthYear']
#     perc_pos_covid_cases = chicago_covid[chicago_covid['MonthYear'] ==
    ↳month_year]['Perc']
#     perc.append(float(perc_pos_covid_cases))

# chicago_combined_dataset_covid['CovidPercPositive'] = perc

# regress_our_model_for_covid(chicago_combined_dataset_covid, 2)
# regress_our_model(chicago_combined_dataset_covid, 2)

```

[29]:

```

# Trying out OLS Regression with multivariate implementation

# import statsmodels.api as sm

# city_data = pd.read_csv('~\Documents\Practice\CSV_LA_city_final_dataset_v2.
    ↳csv')
# normalize_dataset(city_data)
# y = city_data['Scaled_CrimeCount']
# x = city_data.loc[:, ['Gender_ratio', 'Scaled_Unemployment',
    ↳'Scaled_MedianIncomeRate']]

# X = sm.add_constant(x)
# lin_model = sm.OLS(y, X)
# regr_results = lin_model.fit()
# print(regr_results.summary())
# print(list(regr_results.params))

# transformer = PolynomialFeatures(degree=6, include_bias=False)
# x_ = transformer.fit_transform(x)
# x_ = pd.DataFrame(np.array(x_), columns=transformer.get_feature_names_out())
# X = sm.add_constant(x_)
# poly_model = sm.OLS(y, X)
# regr_results = poly_model.fit()
# print(regr_results.summary())

```

```
[30]: # Playing around with the interaction terms

# city_data = pd.read_csv("~/Documents/Practice/chicago_city_final_dataset_v2.
# city_data = add_mental_health_data("CHICAGO", city_data, "~/Documents/
# y = city_data['Cul. Crime count']
# x = city_data.loc[:,['Gender_ratio', 'UnemploymentRate', 'MedianIncomeRate',
# 'WorkforceCount', 'PercentNegativeUsers', 'FavorOfDemocrats',
# 'SearchCountForDepression']]
# transformer = PolynomialFeatures(degree=2, include_bias=False)
# x_ = transformer.fit_transform(x)
# z = transformer.get_feature_names_out()

# # for idx, ele in enumerate(z):
# #     if idx in [3, 7, 10, 11, 12, 13, 14, 15, 17, 18, 19, 20, 22, 23, 24,
# # 25, 28, 30, 31]:
# #         print("{}: {}".format(idx, ele))
# # print(remove_interaction_terms(z, [19, 26, 33, 34]))
# # print(list(z[:10]) + list(z[11:19]) + list(z[20:31]) + list(z[32:]))

# # print(z[:7] + [z[12]] + [z[17]] + [z[21]] + [z[24]] + [z[26]])

# # Removing interaction terms
# new_x_ = []
# interaction_terms = [1, 8, 14, 17, 19, 22, 24, 26, 28, 29, 30, 31, 33]
# for item in x_:
#     temp = list(item[:5]) + list(item[6:10]) + list(item[11:19]) +
# list(item[20:22]) + list(item[23:25]) + list(item[28:31]) + list(item[32:])
#     temp = remove_interaction_terms(item, interaction_terms)
#     new_x_.append(temp)

# # # # new_x_ = pd.DataFrame(np.array(new_x_), columns=list(z[:5]) + list(z[6:
# 10]) + list(z[11:19]) + list(z[20:22]) + list(z[23:25]) + list(z[28:31]) +
# list(z[32:]))
# new_x_ = pd.DataFrame(np.array(new_x_), columns=remove_interaction_terms(z,
# interaction_terms))
# new_x_ = pd.DataFrame(np.array(x_), columns=z)
# X = sm.add_constant(new_x_)
# regr_results = sm.OLS(list(y), X).fit()
# regr_results.summary()
```

```
[31]: # 0: Gender_ratio
# 1: UnemploymentRate
# 2: MedianIncomeRate
# 3: WorkforceCount
```

```

# 4: PercentNegativeUsers
# 5: FavorOfDemocrats
# 6: SearchCountForDepression
# 7: Gender_ratio^2
# 8: Gender_ratio UnemploymentRate
# 9: Gender_ratio MedianIncomeRate
# 10: Gender_ratio WorkforceCount
# 11: Gender_ratio PercentNegativeUsers
# 12: Gender_ratio FavorOfDemocrats
# 13: Gender_ratio SearchCountForDepression
# 14: UnemploymentRate^2
# 15: UnemploymentRate MedianIncomeRate
# 16: UnemploymentRate WorkforceCount
# 17: UnemploymentRate PercentNegativeUsers
# 18: UnemploymentRate FavorOfDemocrats
# 19: UnemploymentRate SearchCountForDepression
# 20: MedianIncomeRate^2
# 21: MedianIncomeRate WorkforceCount
# 22: MedianIncomeRate PercentNegativeUsers
# 23: MedianIncomeRate FavorOfDemocrats
# 24: MedianIncomeRate SearchCountForDepression
# 25: WorkforceCount^2
# 26: WorkforceCount PercentNegativeUsers
# 27: WorkforceCount FavorOfDemocrats
# 28: WorkforceCount SearchCountForDepression
# 29: PercentNegativeUsers^2
# 30: PercentNegativeUsers FavorOfDemocrats
# 31: PercentNegativeUsers SearchCountForDepression
# 32: FavorOfDemocrats^2
# 33: FavorOfDemocrats SearchCountForDepression
# 34: SearchCountForDepression^2

```